Survey of Performance Acceleration Techniques for Network Function Virtualization

In the network function virtualization, this article surveys ways of realizing network functions in a scalable way. It introduces techniques for hardware and software acceleration that allow for removing inherent bottlenecks in the deployment of network functions.

By Leonardo Linguaglossa, Stanislav Lange, Salvatore Pontarelli [ID],
Gábor Rétvári [ID], Dario Rossi [ID], *Senior Member IEEE*, Thomas Zinner [ID],
Roberto Bifulco [ID], Michael Jarschel, and Giuseppe Bianchi

ABSTRACT | The ongoing network softwarization trend holds the promise to revolutionize network infrastructures by making them more flexible, reconfigurable, portable, and more adaptive than ever. Still, the migration from hard-coded/hard-wired network functions toward their software-programmable counterparts comes along with the need for tailored optimizations and acceleration techniques so as to avoid or at least mitigate the throughput/latency performance degradation with respect to fixed function network elements. The contribution of this paper is twofold. First, we provide a comprehensive overview of the host-based network function virtualization (NFV) ecosystem, covering a broad range of techniques, from low-level hardware acceleration and bump-in-the-wire offloading approaches to high-level software acceleration solutions, including the virtualization technique itself. Second, we derive guidelines regarding the design, development, and operation of NFV-based deployments that meet the flexibility and scalability requirements of modern communication networks.

KEYWORDS | Fast packet processing; network function virtualization (NFV); offloading; performance acceleration; virtualization

L. Linguaglossa and **D. Rossi** are with Télécom ParisTech, 75013 Paris, France.
S. Lange is with the Faculty of Mathematics and Computer Science, University of Würzburg, 97074 Würzburg, Germany.
S. Pontarelli is with Axbryd, 00173 Rome, Italy, and also with the National Inter-University Consortium for Telecommunications, 00133 Rome, Italy.
G. Rétvári is with the MTA-BME Information Systems Research Group, H-1117 Budapest.
T. Zinner is with the Department of Telecommunication Systems, Technical University of Berlin, 10623 Berlin, Germany (e-mail: zinner@inet.tu-berlin.de).
R. Bifulco is with NEC Laboratories Europe, 69115 Heidelberg, Germany.
M. Jarschel is with Nokia Bell Labs, 81541 Munich, Germany.
**G. Bianchi** is with the National Inter-University Consortium for Telecommunications, University of Rome Tor Vergata, 00133 Rome, Italy.

I. INTRODUCTION

Current communication and networking use cases such as Industrial Internet, Internet of Things, video streaming, and vehicle-to-everything communication confront operators with numerous challenges. On the one hand, an ever-growing number of devices and services call for scalable systems that can cope with the increasing resource demands. On the other hand, the heterogeneity of demands and communication types and their temporal dynamics require a high degree of adaptability in order to operate the network in an efficient manner while meeting performance criteria.

A possible solution for overcoming the limitations of today's network architectures and addressing these challenges is offered by ecosystems that are built around the paradigms of software-defined networking (SDN)
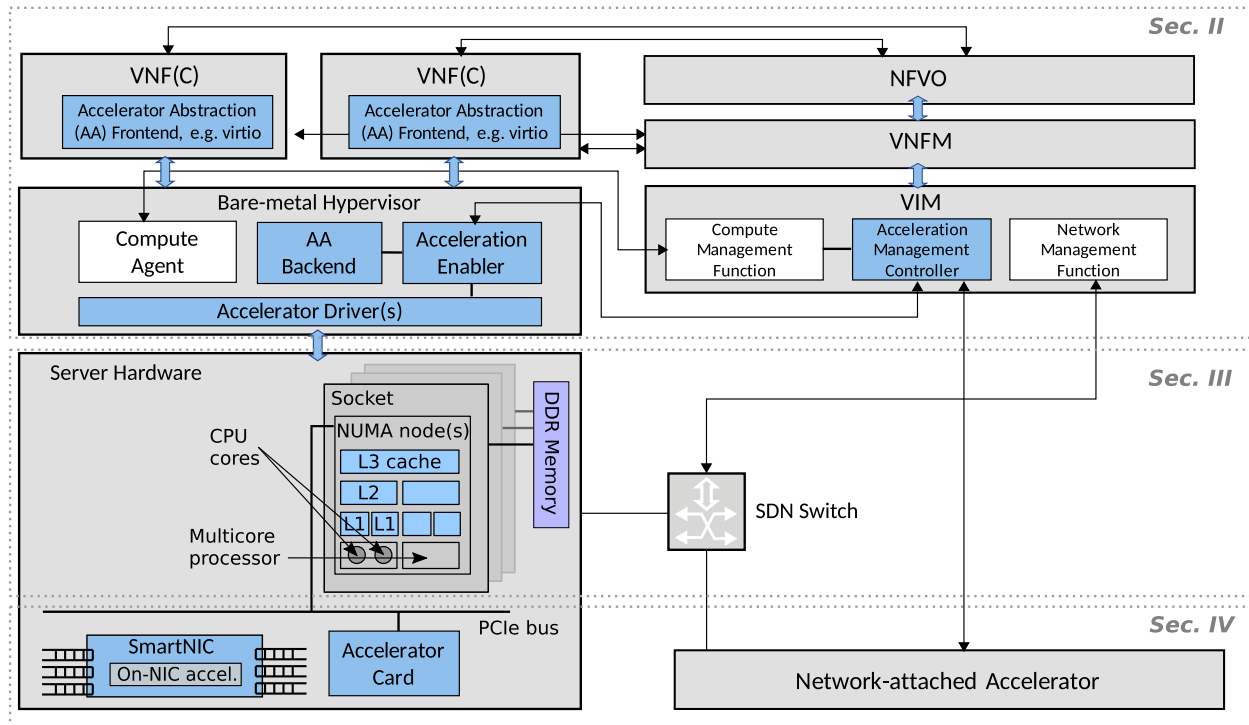
Fig. 1. *Overview of NFV systems considered in this paper. From a top-down perspective, the interaction among the VNFs and the components responsible for the management are explored in Section II. The acceleration techniques can be exploited, via the presence of accelerated drivers, after the deployment of VNFs on server hardware; the most important software acceleration techniques are explored in Section III. Finally, hardware acceleration, which can rely on SmartNICs or external accelerators, is described in Section IV.*

and network function virtualization (NFV). While SDN provides a flexible and programmable infrastructure, NFV replaces dedicated hardware (HW) middleboxes with software instances that can be deployed, scaled, and migrated in a dynamic fashion. However, using software that runs on common-off-the-shelf (COTS) HW to replace special-purpose application-specific integrated circuit (ASIC)-based HW can have a severe performance impact, e.g., in terms of key performance indicators such as throughput and latency, affecting the overall end-to-end application performance. Hence, recent research efforts focus on various acceleration techniques for virtual network functions (VNFs) that address these performance issues at different levels. These include software-based approaches that optimize packets' traversal through the software stack after their reception at the network interface card (NIC), as well as approaches that offload parts of the processing pipeline to programmable hardware. Furthermore, acceleration techniques can be specialized in particular aspects of the processing pipeline and therefore improve the performance regarding resources such as compute, memory, or storage.

One particular problem that arises in such ecosystems is the choice of an appropriate acceleration technique, given a specific combination of network function and virtualization context. There are two main sources of heterogeneity that complicate this task. First, network functions differ significantly with respect to the type and intensity of their resource usage, e.g., compute-heavy transcoding operations as opposed to caches that require a large amount of storage I/O. Second, the virtualization stack, on which the corresponding deployment is built, has a large impact on the availability and applicability of acceleration techniques. Fig. 1 shows a compact representation of the scenario considered in the remainder of this paper. We provide in the figure a reference to the sections devoted to the description of the components of this generic architecture, ranging from pure software components that offer a high degree of programmability to pure hardware devices that have low or no programmability. From a top-down point of view, VNFs are software components built on top of one or more host systems, mainly consisting of commodity PCs. Section II focuses on the ecosystem of VNFs, analyzing the most common network functions and identifying some underlying building blocks. The host system is, in turn, made of several components such as CPU, memory, and storage that are tightly coupled; as an example, the main memory might be split into multiple nonuniform memory access (NUMA) nodes, and the computing unit may consist of several processors, attached to different NUMA nodes. The host system represents a layer on top of which software components, such as VNFs, can be placed. Section III describes the existing pieces of software in scope with the NFV philosophy, as well as the most important software acceleration techniques. Finally, in the lower part of this hierarchy, NICs are responsible

for the actual data exchange among host systems and, therefore, for VNF communication. For some applications, field-programmable gate arrays (FPGAs) or other programmable hardware can be used to offload part of the computation typically provided by the software components. Although all network function may reside in the top of Fig. 1, the virtualized function implemented may have an effect on a different part of the hardware/software hierarchy. Hence, an interesting problem may arise: can all VNFs exploit the same typology of acceleration techniques or there exists some class of VNFs for which a specific acceleration technique is better?

We make the following contributions in this paper. On the one hand, we provide a comprehensive overview of the NFV ecosystem, i.e., functions that are virtualized, virtualization techniques that provide abstractions for accessing physical resources, and orchestration frameworks that integrate all components into the ecosystem. We further decompose VNFs into functional blocks; the characterization of both existing and novel NFs allows to identify potential performance bottlenecks and, consequently, assess the feasibility and expected benefits coming from acceleration mechanisms. In a similar fashion, we survey acceleration techniques. In the case of software acceleration, we provide an overview of mechanisms, their focus and benefits, as well as various tools and frameworks that utilize them. Hardware acceleration mechanisms are subdivided into standard NICs with basic offloading capabilities, SmartNICs that offer FPGA-like programmability, as well as domain-specific programming abstractions for NFs. Based on the outlined survey, we derive the guidelines for the design, development, and operation of NFV ecosystems that can cope with heterogeneous and dynamic workloads in an efficient manner. In addition to the general view of the NFV ecosystem, we choose one exemplary VNF—the next-generation firewall (NGFW)—which leverages a large number of functional blocks and thus represents a good example to illustrate how these guidelines translate to a particular use case.

In the remainder of this paper, we cover the NFV ecosystem (Section II) and survey software (Section III) and hardware (Section IV) acceleration techniques. We next present the guidelines that are extracted from this survey (Section V) and conclude this paper (Section VI).

II. NETWORK FUNCTION VIRTUALIZATION ECOSYSTEM

Numerous areas of networking are targeted by the NFV paradigm [1], [2]. These range from traditional middleboxes [3] such as network address translations (NATs) and firewalls to mobile network functions [4] such as Serving Gateway (SGW) or Packet Gateway and mobile management entities, as well as further extending to customer premises equipment and media transcoders. In this section, we provide a structured overview of the NFV ecosystem in order to cope with the resulting heterogeneity in terms of requirements for an efficient operation of individual functions as well as management of their coexistence on a shared physical substrate.

First, we break down VNFs into common functional building blocks (Section II-A) to identify consumed resources, possible performance bottlenecks, and therefore opportunities for acceleration. This modular approach allows for a straightforward classification of new functions and their building blocks. Furthermore, guidelines and insights for functions that are composed of the same building blocks can be transferred between functions. Second, we discuss different virtualization options (Section II-B) and how they are used to achieve sharing and isolation of network, compute, storage, and memory resources. Finally, we provide a brief overview of options regarding management and orchestration (MANO) mechanisms (Section II-C) alongside their compatibility with the aforementioned virtualization mechanisms.

A. Set of Network Functions and Building Blocks

As outlined earlier, there are numerous different VNF types and areas of use. Specialized surveys and articles provide information regarding individual functions and their details [1], [2], [4], [5]. In contrast, we first identify common building blocks that can be used to compose and classify a particular function. We then demonstrate how this classification and decomposition can be applied to the NGFW which features a large number of these building blocks, and later present guidelines regarding acceleration options for the NGFW.

1) Individual Functions and Blocks: Each building block is characterized by a certain degree of resource utilization (e.g., compute, memory, and storage) as well as the position within a function at which the corresponding action is usually performed. For several VNFs, Table 1 displays their decomposition into building blocks alongside the each blocks' characteristics, which are outlined next.

a) Header read: Basically, all network functions that are considered in this paper need to parse packet headers as a part of their processing and decision-making process. Headers are typically compact and well defined, often in the binary or type length value format, for which this operation is neither compute- nor memory-intensive and requires no access to storage.

b) Header write: In order to perform the tasks of an NAT function or enable tunneling for virtual private network (VPN) or even for simple IP forwarding, NFs also need to modify the content of packet headers. Although the complexity of such a write access depends on the particular type of operation and incurs a larger CPU and memory usage when compared with a read-only header access, it still amounts to a negligible factor on modern systems.

c) Table lookup: Often, to make decisions regarding network actions (e.g., like the appropriate forwarding of an IP packet or an Ethernet frame), a lookup based on some packet properties (e.g., the IP or Ethernet destination

Table 1 Network Functions, the Building Blocks They Are Composed of, and Characteristics of the Building Blocks With Respect to Their Resource Usage. BS: Base Station. BBU: Baseband Unit. DPI: Deep Packet Inspection. IPS: Intrusion Prevention System. FW: Firewall. HTTP HE: HTTP Header Enrichment. IDS: Intrusion Detection System. NAPT: Network Address Port Translation. NGFW: Next Generation Firewall. TCP Opt: TCP Optimizer. VPN: Virtual Private Network

Network Function	Headers		Table	Payload		Flows		Session	Misc	
	(1) Header Read	(2) Header Write	(3) Table Lookup	(4) Payload Inspection	(5) Payload Mod.	(6) Flow State Tracking	(7) Flow Actions	(8) Session Mgmt.	(9) Disk Access	(10) Signal Proc.
BS / BBU										✓
Cache	✓	✓	✓	✓			✓		✓	
DPI, IPS	✓		✓	✓		✓	✓	✓		
FW, Load Balancer	✓		✓			✓	✓			
HTTP HE	✓	✓	✓			✓		✓		
IDS	✓		✓	✓		✓		✓		
Monitor	✓		✓	✓		✓		✓	✓	
NA(P)T	✓	✓	✓			✓				
NGFW	✓	✓	✓	✓	✓	✓	✓	✓		
Proxy	✓	✓	✓	✓		✓		✓		
Shaper	✓	✓	✓			✓	✓			
TCP Opt, VPN	✓	✓	✓	✓	✓	✓	✓			
Transcoder	✓		✓	✓	✓	✓		✓		
Characteristic	(1) Header Read	(2) Header Write	(3) Table Lookup	(4) Payload Inspection	(5) Payload Mod.	(6) Flow State Tracking	(7) Flow Actions	(8) Session Mgmt.	(9) Disk Access	(10) Signal Proc.
Compute	ϵ	ϵ	+	+ / ++	++	+ / ++	ϵ	++	+	++
Memory	ϵ	ϵ	+	+ / ++	+ / ++	+ / ++	ϵ	++	+	+ / ++
Storage	-	-	-	-	-	+	-	+	++	-
Position	Start	End	Start	Start	Middle	Middle	End	Middle	Middle	Middle

"-", "ϵ", "+", and "++" represent no, low, medium, and high usage of the corresponding resource, respectively.

address) in a specific table (e.g., typically implemented as tries in the case of IP forwarding or as hash-based data structures for Ethernet) is necessary. These tables are typically kept in memory so that no additional storage is required. Hence, CPU usage and memory I/O are higher, particularly in the case of larger forwarding tables.

d) Payload inspection: More sophisticated actions, such as deep packet inspection (DPI) that is performed by an intrusion detection system such as Bro [6] or Snort [7] or by an NGFW, take into account packet payload to identify malicious content and applications. Depending on the complexity of the performed analysis, this block can be both highly compute- and memory-intensive.

e) Payload modification: Similar to the case of header read and write operations, payload modification requires additional CPU cycles for copying packet contents to memory. Furthermore, functions, such as transmission control protocol (TCP) optimizers and VPNs, perform compute-intensive tasks related to compression and encryption, respectively.

f) Flow state tracking: Although they perform actions on a per-packet basis, several functions need to maintain per-flow state information. This information can range from simple header characteristics in an NAT lookup table, to data that is more demanding in terms of CPU and memory resources such as a short-term history of a flow's packets for tasks such as monitoring and intrusion detection [6], [8] or resource sharing for concurrent flows [9].

g) Flow actions: Furthermore, functions, such as firewalls and load balancers, perform actions on a per-flow basis. These include flow termination or redirection, are usually performed at the end of a function, and consume a few resources.

h) Session management: Several functions, such as intrusion prevention systems (IPSs), need information on multiple flows in order to reconstruct and analyze the corresponding application sessions and persistently store data that are relevant for accounting purposes. These tasks can incur an even higher demand for CPU and memory resources than their per-flow counterpart. Similar to the case of flow state tracking, some information might be logged to disk.

i) Disk access: A subset of NFs heavily depends on read and write access to persistent storage (e.g., nDPI [10]). These NFs include caches as well as network probes that capture and log traffic traces at either packet level or flow level.

j) Signal processing: Finally, CPU-intensive signal processing tasks, such as the fast Fourier transform, need to be performed by the cloud radio access network functions such as the baseband unit in a virtualized base station.

As mentioned earlier, building blocks can not only be mapped to the consumption of a particular resource but also to a position within a function. This factor is relevant when considering to accelerate a building block since its position can affect the feasibility and performance gain of the acceleration. While we acknowledge that the position of a building block can vary from function to function, we observe certain patterns, e.g., reading headers and payloads are most likely performed at the beginning of a processing pipeline, whereas flow-level actions are usually performed at the end.

2) Next-Generation Firewall: In order to illustrate the characterization of a network function by means of its building blocks, we consider the NGFW that is outlined in [11]. In addition to standard firewall features, the NGFW is also capable of performing DPI and encompasses IPS as well as Secure Socket Layer VPN functions and man-in-the-middle proxying capabilities. We deliberately choose this VNF due to its wide range of functionalities and the resulting complexity of characterizing it in its entirety. From the capabilities of the NGFW, we can identify its building blocks and requirements; after reading the header of a received packet (1), the NGFW needs to check the packet's flow and/or session membership (3, 6, 8). Furthermore, a scan of the packet's payload might be necessary (4) before deciding whether to forward, redirect, or drop the packet (7). Finally, packet headers need to be rewritten (2), sometimes also with payload modification (5), e.g., in the case of proxying or for encryption in the case of VPN tunneling.

B. Virtualization Techniques

In order to fully reap the flexibility and scalability benefits of the NFV paradigm, the underlying infrastructure needs to expose access to a shared resource pool. Simultaneously, it needs to maintain isolation between different tenants and dynamically adapt the resource allocation to VNF instances for scalability and efficiency. In this context, conventional host resources, such as compute, memory, and storage, are virtualized alongside networking resources, i.e., links between nodes with delay and bandwidth guarantees. In the following, we provide an overview of virtualization mechanisms from these two categories and refer to them as: 1) hardware virtualization and 2) network virtualization (NV), respectively. The choice of options from these two categories results in different tradeoffs regarding costs, performance, and flexibility.

1) Hardware Virtualization: The main goal in the context of hardware virtualization consists of providing a functional equivalent of the original hardware to the VNF while limiting performance degradation. This is achieved by the virtualization layer that resides between the physical and the virtualized infrastructure. Options regarding its implementation range from different hypervisors to container-based approaches, which offer different degrees of resource and isolation, security, and performance [2], [12].

Hypervisors can reside at different locations in the system. On the one hand, hypervisors, such as Xen [13] and ESXi [14], run directly on the physical machine and directly handle hardware access as well as interrupts. On the other hand, hypervisors, such as the kernel-based virtual machine (KVM) [15], require the presence of kernel modules in the host OS and kernel processes are responsible for handling events and interrupts. Due to the different numbers of intermediate abstraction layers, implementation details, and the heterogeneity of possible use cases, there is no single hypervisor that performs best in all scenarios and therefore should be chosen carefully to match the particular requirements [16]. Furthermore, high-speed I/O virtualization plays a crucial role in meeting performance requirements and maximizing flexibility when VNFs are run inside virtual machines (VMs) [17]. To this end, drivers, such as virtio [18] and ptnet [19], provide a common device model for accessing virtual interfaces.

Recently, container-based technologies, such as Docker [20], have gained popularity, which instead share the kernel of the host OS and provide isolation through segregated process namespace (anciently known as Berkeley Software Distribution (BSD) jails [21]). On the one hand, containers avoid the overhead of providing an entire guest OS, enabling a flexible and lightweight deployment. In particular, network ports or entire interfaces can be directly assigned to a container without incurring an overhead as in the case of hypervisor-based solutions. On the other hand, sharing the kernel also raises concerns regarding security and (performance) isolation between instances. Finally, worth mentioning are also unikernel [22] as well as serverless [23], [24] technologies for the virtualization of network functions, which sits at the opposite side of the spectrum as the former targets full flexibility with lightweight monolithic kernels that run on a hypervisor such as Xen, whereas the latter targets performance by trading it for a more constrained environment offering a set of application program interfaces (APIs). All these technologies are still competing, and there is no single candidate that fits all use cases, which makes the ecosystem very rich.

2) Network Virtualization: In a similar fashion to hardware virtualization, NV accomplishes an abstraction of network resources and therefore acts as an enabler for fully programmable networks whose entire stack can be provisioned in a flexible manner. The survey in [25] provides an overview of NV approaches that are categorized with respect to several characteristics. These include

the network technology that is virtualized, the layer in the network stack at which virtualization is performed, as well as the granularity of virtualized entities. Depending on the envisioned use case, operators need to decide on an NV technology that satisfies the corresponding requirements.

An additional step in the evolution of NV consists of combining the NV idea of sharing a physical substrate among several stakeholders with the SDN paradigm. By introducing an SDN hypervisor, it is possible to virtualize SDN-based networks and therefore allow multiple tenants to deploy their own SDN-based network including their own controllers while operating on one common physical infrastructure. A survey of SDN hypervisors is available in [26]. The main distinctions between different hypervisors include their architecture, the network attributes that are abstracted, as well as the capabilities in terms of isolation between slices.

This paper leverages the technologies described in [26] and clearly build upon [25]; specifically, we go one step further with respect to [25] by understanding how functions can be implemented and, especially, accelerated with either software, hardware, or hybrid approaches.

C. Integration Into the Orchestration Context

Given the abstractions and flexibility of the outlined virtualization mechanisms, MANO techniques can attach to interfaces that are defined in frameworks such as ETSI MANO [27] and adapt the resource allocation of individual VNF types and instances in a dynamic fashion. Additionally, knowledge regarding VNF's resource requirements and performance levels given a specific resource enables scaling VNFs according to observed demands. An overview of efforts from research projects, standardization bodies, and both open-source and commercial implementations is provided in [28].

A concrete example regarding the integration of hardware acceleration mechanisms into the MANO context is outlined in [11]. In the presented use case, a virtual infrastructure manager enables the elastic deployment of an IPSec tunnel termination VNF by combining two types of information. On the one hand, it has knowledge of nodes' CPU cores, available instruction sets, and hardware acceleration capabilities. On the other hand, the VNF descriptor provides the performance levels in terms of possible bandwidth, given different numbers of CPU cores, instruction sets, and hardware acceleration.

Whereas orchestration is out of the scope of this paper, [11] and [28] testify that effort is ongoing into making it possible to harness acceleration techniques in a seamless way. It is thus important to understand which of the software or hardware acceleration is suitable for VNFs and more generally to outline a pattern for the design and implementation of new VNFs that can make the most out of the available acceleration techniques.

III. SOFTWARE ACCELERATION TECHNIQUES

To bridge the performance gap between specialized hardware middleboxes and software instances that run on COTS hardware while retaining the flexibility of the latter, the past two decades have seen tremendous advances in software acceleration techniques (Section III-A). This sparked a quick evolution of pure-software user-space network stacks and NFV frameworks, which attains multi-gigabit performance on COTS platforms (Section III-B).

A. Software Acceleration

We review the most important software acceleration techniques in the context of NFV. In the following discussion, we make a distinction between pure software acceleration techniques, program code optimization approaches affecting the parts of the network stack, including firmware, driver code, operating system code, and the user-facing network infrastructure, which can be directly modified by a user or a vendor without having to upgrade or modify the underlying hardware, and hardware-assisted acceleration techniques that require explicit support in hardware. Table 2 provides a compact summary of the techniques and their benefits.

1) Pure Software Acceleration: Polling techniques, as opposed to an interrupt-driven design, have been used for quite some time to speed up NIC drivers. In the interrupt-driven mode, the NIC driver is managed by the operating system in the kernel-space or a user-space application via kernel bypass and is responsible for pushing received packets to main memory and to interrupt the CPU to signal the availability of new packets to process. The CPU then enters a special interrupt-handling routine and starts VNF processing. However, the corresponding context switches may cause nonnegligible overhead at high loads and the interrupts may saturate the CPU. Interrupt-mitigation techniques have a relatively long history; one such technique is represented by interrupt coalescing [29], which waits for a fixed timer for further packet reception before raising an interrupt. Yet, a more efficient technique to avoid overloading the CPU at very high packet rates is switching to polling mode, whereby the CPU periodically checks for pending packets (pull model) without the need for the NIC to raise an interrupt (push model). Polling the NIC in a tight loop, however, results in 100% CPU usage regardless of the traffic load. Support for poll-mode drivers has been introduced in BSD at least since 2001 [30], and interrupt coalescing is also available in the Linux kernel [31].

I/O batching is another technique that is commonly used in poll-mode NIC drivers to mitigate interrupt pressure at high load. In the batched mode, the NIC does not receive or transmit packets individually but operates by aggregating several packets into a batch. Upon reception of a packet, the NIC writes it into a hardware queue, and then, the whole batch is pulled by the CPU in

Table 2 Benefits of Software Acceleration Techniques. ZC: Zero Copy. MP: Mempools. HP: Hugepages. PF: Prefetching. CA: Cache Alignment. LFMT: Lock-Free Multithreading. LT: Lightweight Threads. ML: Multiloop. BP: Branch Prediction. RSS: Receive-Side Scaling. FH: Flow Hash. SIMD: Single Instruction Multiple Data. DDIO: Direct Data I/O. SR-IOV: Single-Root Input/Output Virtualization

	Poll	I/O Batch	Memory					Compute Batch	Threading		Coding		NIC-support		CPU-support		
			ZC	MP	HP	PF	CA		LFMT	LT	ML	BP	RSS	FH	SIMD	DDIO	SR-IOV
Reduce memory access			✓		✓	✓	✓									✓	
Optimize memory allocation				✓	✓		✓										
Share overhead of processing								✓			✓						
Reduce interrupt pressure	✓	✓															
Horizontal scaling									✓	✓			✓				✓
Exploit CPU cache locality					✓	✓			✓							✓	
Reduce CPU context switches	✓	✓							✓	✓							
Fill CPU pipeline								✓			✓	✓			✓		
Exploit HW computation														✓	✓	✓	✓
Simplify thread scheduling	✓									✓							

a single chunk. In the reverse direction, the CPU may wait for multiple packets to accumulate in the transmit queue before performing an I/O operation. Accordingly, I/O batching overcomes the bottleneck that exists between the NIC and the main memory by amortizing I/O costs over multiple packets. The price to pay is an increased per-packet latency, due to packets having to wait in the queues until a complete batch is formed. First used in PacketShader [32] and DoubleClick [33], I/O batching has become widely adopted by all major high-speed packet processing frameworks to today [34]–[36].

Optimizing memory management is another technique that may significantly improve the performance of VNFs. Upon finalizing packet reception, the kernel needs to make the received data available for a user-space application to consume it. This operation may be an additional bottleneck in traditional network stacks, since the memory region where packets can be written by the NIC via direct memory access (DMA) may not overlap with the user-space memory, incurring a memory copy overhead. Several techniques exist to overcome this limitation. A zero-copy network receive/transmit code path [35]–[37] mitigates costly memory operations by mapping DMA regions into user-space memory and passing only lightweight packet descriptors between the kernel and the application; such descriptors may contain a pointer to the packet buffer and some additional metadata.

Memory management, i.e., the allocation of new packet buffers or the freeing of unused ones, may also be a source of significant overhead. To avoid memory management in the fast packet processing path, modern network stacks typically use preallocated packet buffers. A set of memory pools is created during start-up and this memory area is never freed; such mempools are usually organized into ring buffers where packets are written to, and read from, in a sequential fashion. Mempools are usually allocated from hugepages whose size typically ranges from 2 MB to 1 GB to prevent misses in the translation lookaside buffer. Care must be taken, however, to place mempools to the NUMA node where the majority of memory accesses will be made during runtime; otherwise, performance may incur the penalty of accessing faraway memory across NUMA nodes. Network programmers should also take care of adopting cache-friendly data structures; for instance, hash-table buckets should always be sufficiently aligned to cache lines and should occupy as a few cache lines as possible in order to maximize CPU cache hit rate, and compressed data structures may be used to reduce the overall memory footprint of performance-sensitive data [38]. Finally, prefetching data ahead of time from main memory may substantially contribute to increased performance, by avoiding CPU stalls (i.e., the execution of CPU instructions is blocked due to memory access).

Compute batching is a technique that can be used to enhance the CPU pipeline performance [34], [39]. Traditional network engines typically process network traffic on a per-packet basis; network functions in the processing path are sequentially applied to a packet until the final forwarding decision is made. With compute batching, the notion of I/O batching is extended to VNF processing, in which network functions are implemented from the outset to work on entire bursts of packets rather than on a single packet. Similar to I/O batching, compute batching helps mitigating the overhead of invoking VNFs

(e.g., context switches and stack initialization) as well as providing additional computational benefits. First, compute batching optimizes the use of the first-level CPU instruction cache; when a packet batch enters a network function, the corresponding code is fetched into the L1 instruction cache upon processing the first packet, and the instruction cache remains warm during processing the rest of the packet burst. Furthermore, using compute batching allows to take advantage of data-level parallelism primitives available in modern CPUs (see SIMD in the following). Finally, compute batching also allows better use of the CPU processing pipeline, as it facilitates the CPU pipeline to be consistently full, which allows the CPU to execute multiple instructions within the same clock cycle [34].

Exploiting CPU parallelism has become imperative, given that to manage CPU power dissipation, processor makers started favoring multicore chip designs. Thus, network applications must be written to exploit a massively multithreaded design. Packet processing presents a workload that is ideally suited for parallelization, in which each packet or packet batch may be serviced by a separate execution thread, but care must be taken to avoid packet reordering and, especially, to preserve flow coherence (in both the forward and backward flow directions). To maximize multicore scalability, it is imperative to avoid synchronization between the threads as much as possible, e.g., by preferring lightweight mutual exclusion primitives such as read-copy-update over mutexes and semaphores, and leveraging atomic instructions, lock-free data structures, and per-CPU variables. Lock-free parallelism is tightly coupled with the availability of multiple hardware queues, since different CPUs can be attached to different hardware queues [see receive-side scaling (RSS) in the following] to operate on different traffic subsets. More recently, frameworks, such as Data Plane Development Kit (DPDK) [35], started complementing the fully fledged preemptive multitasking offered by the underlaying OS, with the introduction of a simpler cooperative multitasking model. With the standard run-to-completion execution, each packet is assigned to a separate thread that performs complete VNF processing until the final forwarding decision. In a cooperative environment, a thread can implement a subset of functionalities and can decide when to pass the packet to another thread for further processing. This was made possible by the introduction of lightweight threads, making the VNF scheduler logic much simpler and faster.

Programming best practices also becomes imperative for network application developers, as some coding techniques are known to help compilers to accelerate packet processing. One such best practice is called multiloop programming, an evolution of the conventional loop unrolling practice by explicitly writing the program control flow in a way as to handle packets in groups of two or four in each iteration (dual loop or quad loop, respectively). Loop unrolling has two main goals. The first goal is to reduce the number of "jump" instructions (that can introduce idle slots in the CPU pipelines). The second goal is to carefully prefetch new packets while operating on already prefetched packets; this coding practice helps the compiler to organize the code in a way as to maximize the usage of the CPU pipeline [34]. In addition, explicitly inlining frequently accessed code may allow to save on runtime function call overhead, at the cost of a slightly increased program text size [40]. Finally, using workload-related insights, a network programmer may manually annotate the branches of conditional program code more likely to execute during runtime, this way improving the CPU's branch prediction success rate. Using such programmer-provided annotations, the compiler can place the likely execution branch right after the conditional expression and let the CPU automatically fetch the corresponding code into the CPU pipeline; in contrast, the CPU pipeline must be tediously invalidated and repopulated with the correct branch code in case of any mispredicted branch, possibly leading to a significant performance penalty.

2) Hardware-Supported Functions in Software: An important class of acceleration techniques is constituted by hardware-assisted functions, whereby the hardware exposes certain functionalities that can be used to speed up the execution of software. We distinguish two categories in this context, depending on whether the assistance is offered by the NIC or by the CPU, which are reported in the rightmost part of Table 2.

NIC-assisted acceleration techniques range from virtualization support and direct DMA to NIC-driven parallel packet dispatching. Modern NICs contain a fast packet parser to compute flow-level hashes in hardware, maintain multiple hardware packet queues called RSS, and expose packet counters in registers. Access to this hardware functionality is typically implemented in the low-level NIC drivers. The use of register-backed packet counters reduces memory overhead, while RSS is instrumental to multithreading process, as the packet RSS hash may be used by the NIC to dispatch packets to different CPU cores, in order to leverage flow-level parallelism and avoid packet reordering. RSS hashes ensure that packets of a single transport-level connection (and, depending on the RSS seed, of both directions of the connection) will always be scheduled to the same CPU, which also enforces locality of data structures usage. Furthermore, newer NICs can take advantage of Data Direct I/O, which allows packets to be transferred directly into the last-level CPU caches instead of into the main memory, preventing a costly cache miss when the CPU starts processing the packet. Single-root input/output virtualization (SR-IOV) in addition lets the NIC arbitrate received packets to the correct VNF without the explicit involvement of a hypervisor switch.

CPU-assisted acceleration techniques, on the other hand, leverage the features of modern CPUs to speed up network applications. Most modern CPUs support low-level data parallelism through an advanced single-instruction

multiple data (SIMD) instruction set. SIMD operations allow to execute the same instruction on multiple data instances at the same time, which greatly benefits vector-based workloads such as batch packet processing. In addition, the latest CPU chipsets include built-in CPU virtualization support, exposing a single CPU as multiple virtual CPUs to different VNFs, hyperthreading to multiply the number of CPU pipelines for better multicore scalability, and multiple memory channels that allow applications to load-balance memory accesses.

B. Ecosystem of Software Stacks

The previous techniques are leveraged by a number of tools that we overview in this section. Roughly, software stacks can be categorized into three main branches, depending on whether they are: 1) low-level building blocks; 2) specialized pieces of software implementing a single VNF; or 3) full-blown frameworks for network function composition. In the following, we discuss the representatives from each category and provide a visual overview of the resulting ecosystem of stacks in Fig. 2.

1) Low-Level Building Blocks: Low-level high-speed packet processing libraries may be implemented in the kernel space or in the user space using kernel bypass. Kernel-space solutions have full access to the resources of the machine but must be accessed via standard system calls (the kernel performs controls to grant fault tolerance and isolation with respect to other processes). On the contrary, user-space approaches bypass this step via specific libraries, thus avoiding the overhead of the kernel at the cost of a reduced isolation. Particularly, while kernel-based networking has progressed at a relatively low pace, user-space data-plane libraries have seen tremendous advance within the last decade, with a flourishing ecosystem of which [35]–[37] are representative examples. A performance comparison of these frameworks for packet forwarding is available in [39].

Netmap [36] is a framework for packet I/O that provides a high-speed interface between the user-space applications and the NIC (with assistance of a kernel-level module). Depending on the NIC utilized, netmap may leverage hardware capabilities (such as RSS queues), and it allows to reach a 10-Gb/s rate with a single low-tier CPU (900-MHz CPU clock rate). Being part of FreeBSD and also available on Linux, netmap is well-known in both the research and industrial communities and is typically used to build packet-sniffing or traffic-generator applications, L2/L3 software routers, and even firewalls. The Intel DPDK [35] is another framework for high-speed packet processing with a large user community. In contrast to netmap, the DPDK provides a full kernel bypass, exporting the NIC driver logic as well into user space and exposing a rich NIC abstraction to applications; registers, counters, and hardware-assisted functionality may all be accessed by the programmer using a convenient C language API. PF_RING_ZC [37] is another
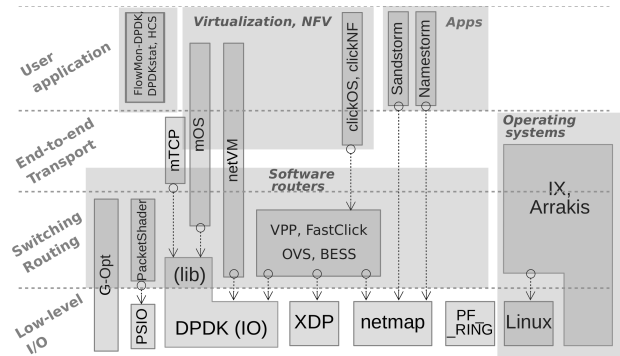
Fig. 2. *Ecosystem of software stacks for high-speed network function composition.*

multi-10-Gb framework targeting interprocess and inter-VM communications. PF_RING_ZC NIC drivers may work both as a regular kernel driver or can be switched into kernel-bypass mode, allowing the user to select which packets are received in the kernel-bypass mode and which ones get processed by the standard Linux kernel network stack.

Recently, the above-mentioned features used by kernel-bypass stacks have started being introduced in the Linux kernel, notably by AF_PACKET v4 [41] and the eXpress Data Path (XDP) [42] projects. XDP adopts many of the aforementioned acceleration techniques, such as batched I/O, multiple HW queues, and zero copy, on top of a virtual machine that allows the user to inject custom packet processing programs into the network stack written in the C-like enhanced Berkeley Packet Filter (eBPF) language. Applications can also take advantage of the MSG_ZEROCOPY feature available in recent Linux kernels that enables zero-copy transmission of data, and efforts are ongoing to address the receive code path as well [43]. Given the recent appearance of these functions in the kernel, a comprehensive comparison of high-speed kernel and kernel-bypass stacks is still lacking so far.

2) Monolithic Network Functions: One step above in the ecosystem is represented the realization of useful network functions based on the low-level building blocks; a set of monolithic network functions have started to appear that address one specific use case, such as forwarding, NAT, intrusion and detection, and transcoding. These functions are custom stand-alone software components that either (more rarely) reimplement a custom user-space network stack making use of the acceleration techniques early introduced or (more frequently) build their VNF-specific code on top of an existing low-level building block, such as DPDK or netmap.

One such specialized VNF is PacketShader [32], and a high-speed packet I/O framework targeting packet processing functionality is implemented, such as IPv4/IPv6 forwarding and OpenFlow switching. PacketShader greatly benefits from GPU acceleration. Another tool is G-opt [44], a software router which targets memory

access optimization to speed up packet processing. The hierarchical content store [45] targets the problem of line-rate access to a router cache in information-centric networking [46].

Sandstorm and Namestorm [47] are specialized user-space network stacks building over netmap to implement Web server and DNS server functionalities, respectively. DPDKStat [8] and FlowMon-DPDK [48] are, as the name implies, DPDK-based tools for traffic monitoring; the former is a CPU-hungry tool providing features such as DPI, full payload analysis, flow reconstruction, and exports hundreds of advanced flow-level statistics, and the latter targets very simple flow tracking with minimal CPU resources. mTCP [49] is a user-space TCP framework targeting high-speed flow-level packet processing, designed to exploit multicore systems, and it can work on top of Linux as well as PacketShader I/O, DPDK, or netmap.

Finally, Open vSwitch (OVS) [50] and ESwitch [40] are examples of fully fledged high-performance OpenFlow software switches that come with some network functions built-in (e.g., NAT with connection tracking). OVS, in particular, is inspired by the SDN paradigm of the separation between control and data plane and, as such, works with a set of match/action rules, similar to a regular SDN switch, and it can be deployed both over a stock Linux kernel using kernel-space networking and also over a fast DPDK datapath.

3) Modular Network Function Composition Frameworks:
One last relevant class of the software stack ecosystem is represented by frameworks whose main aim is generality and feature richness; in contrast to low-level blocks or monolithic functions, these frameworks provide a set of already available VNFs, as well as an environment which simplifies the development of new VNF and, especially, their composition into complex and dynamic service chains.

In this category, we find software routers, such as vector packet processor (VPP) [34], Click [51] and its high-speed variant FastClick [39], and Berkeley Extensible Software Switch (BESS) [52]. These frameworks allow users to configure, manage, and deploy flexible service chains by developing and customizing monolithic network functions such as ACL and NAT as elementary building blocks and then arranging these network functions into a dataflow graph. At high level, the difference is that VPP is aimed at typical L2–L4 workloads and, as such, the dataflow graph is fixed, whereas Click and BESS allow arbitrary processing graphs, thus focusing more on flexibility and reconfigurability.

While the previous frameworks delegate the majority of the work for developing and arranging VNFs to the user, other frameworks offer a "cloudlike" environment for the flexible composition of network functions encapsulated as stand-alone virtual machines or containers. Examples are ClickOS [53], ClickNF [54], mOS [55], OpenStack [56], and NetVM [57]. ClickOS [53] provides tiny virtual machines (unikernels) that implement the Click abstraction for providing resource-efficient and high-performance virtual machines. ClickNF [54] and mOS [55], respectively, build over Click and mTCP to offer the possibility to design, prototype, and deploy custom middleboxes on commodity hardware. In addition, OpenStack [56] and NetVM [57] provide telco cloud services to deploy customizable data-plane functionality as virtual machines, targeting a high-speed inter-VM communication paradigm.

Finally, worth mentioning are IX [58] and Arrakis [59]. Rather than targeting composition of VNF through a dataflow graph abstraction or in a cloudlike NFV infrastructure, these frameworks offer a new operating system that is explicitly tailored for network function deployment. In particular, both IX and Arrakis reimplement basic operating system functionality, including I/O, memory management, and system calls, to provide a protected environment in which high-speed network functions can be deployed as simple network applications.

Given the abundance of tools for VNF deployment with similar spirit and capabilities, making a selection of the most appropriate candidate is not an easy task. Particularly, a very few works exist that aim at independently and directly benchmarking these frameworks in a scientifically sound way. Further work beyond [60], which compares OVS, SR-IOV, and VPP, is thus necessary to provide a complete picture of the ecosystem.

IV. HARDWARE OFFLOADING

With the end of Dennard's scaling and Moore's law, general-purpose CPUs cannot increase their processing power at the same pace at which per-interface network speeds grow [61]. As a matter of fact, NICs already offload several functions that should be otherwise provided by a system's CPUs. For example, commodity NICs, such as those described in [62], [63], can offload checksum calculation and verification, packet filtering, and functions related to specific protocols such as VxLAN encapsulation or TCP segmentation. Such operations help the system in handling traffic rates in excess of 40 Gb/s per NIC port, saving the CPU for further computations.

In line with this trend and in order to handle a larger number of potentially complex functions, the adoption of programmable network function accelerators is becoming widespread.[1] Generally, these accelerators are integrated in the NIC, hence called a SmartNIC [65]–[67], and deployed as bump-in-the-wire, i.e., they apply their processing either before or after the processing performed by the system's CPUs.

Current SmartNICs' programmable engines are usually provided with proprietary software that offloads some well-known protocols and functions, whose programmability is not exposed to third-party users. This allows to

---

[1]Such a need has also been recently recognized by standard organizations such as ETSI ISG NFV, which is defining hardware accelerators' capabilities in their specifications for NFV platforms [64].

quickly adapt the SmartNIC functions to emerging needs, without waiting for the long implementation cycles experienced with traditional NIC hardware [68], but it is a process still largely driven by specific vendor's interests. For example, it is common for SmartNICs to provide only well-known functions, such as firewall offloading, e.g., integrated with Linux's iptables, or offloading of virtual switches, e.g., OVS [50].

However, the offloading of fixed functions hardly matches the requirements of modern dynamic and fast changing applications [69]–[72]. As a result, recently, vendors started to increasingly open their programmable engines and provide users with the ability to customize the SmartNIC's software. As such, it becomes of particular interest to study how SmartNICs can be leveraged to become programmable network function accelerators and what are the related opportunities and issues.

A SmartNIC [73] typically adopts specialized network processors (NPs) or FPGAs as programmable engines. Despite obvious considerations regarding costs, with the NICs being usually cheaper than SmartNICs, the latter also have important implications in terms of ease of use, application development time, and performance guarantees, to name a few. In particular, while a Smart-NIC gives great flexibility to introduce new functions enhancing system performance, it comes with the complexity of programming their engines, which includes phases of design, debugging, verification, and optimization. To overcome this design complexity, SmartNICs often offer higher level, domain-specific programming abstractions for their programmable engines [74], [75]. Such abstractions implement a point in a conceptual design space that ranges from fixed functions, on the one side, to full programmability, on the other, with any specific solution trading in its way programming simplicity with expressiveness.

In this sense, one of the open research questions for SmartNICs is about which abstraction should be adopted in order to provide ease of use, flexibility, and performance at the same time. In the remainder of this section, we first describe the primitives usually offloaded to commodity NICs (Section IV-A), then give a short survey of the Smart-NIC architectures (Section IV-B), and finally describe the currently available programming abstractions for Smart-NICs' offloading engines (Section IV-C).

Here, we want to underline that there are complex tasks, such as DPI, which are appealing candidates for hardware acceleration but that are not covered in this paper. For example, DPI is a candidate for hardware acceleration due to the cost of payload inspection and to the complexity of pattern matching operations. In fact, DPI acceleration has been extensively considered both using different targets such as GPU [76] and FPGA [77] and using different algorithmic approaches [78]–[80]. However, due to the specific nature and complexity of such functions, the corresponding acceleration techniques require an *ad hoc* and extensive description that is outside the scope of this paper.

We instead refer the interested reader to [81] for further details.

A. Offloading in Commodity NICs

Here, we describe a set of commonly used primitives offloaded to NICs and discuss to what extent these offloading procedures can be made programmable in order to satisfy the flexibility required by NFV applications. A set of widely offloaded primitives is listed next; clearly, commodity NICs offer a set of well-defined fixed functions targeting existing widespread protocols.

1) CRC/Checksum Computation: The widespread use of some checksum computation, such as the 1s complements sum used by TCP/IP protocols, makes it suitable for hardware offload. In transmission, the NIC directly computes the checksum value defined by the specific protocol and sets the corresponding field in the packet before sending it out; in the receive mode, the NIC checks the checksum's correctness and discards the packet if the checksum is wrong. Almost all commercial NICs provide TCP/IP checksum offloading. Since the Stream Control Transport Protocol (SCTP) checksum algorithm is different from TCP/IP, only a few NICs support SCTP checksum offloading. This is clearly related to the smaller adoption of this protocol with respect to TCP. It is worth noting that checksum computation could be easily performed by current generation processors, also exploiting specific instruction set extensions to speed up the computation. This limits the gain between hardware and software implementations. However, the cost of implementing such operations in hardware is marginal, thus making its offloading appealing. Furthermore, offloading the checksum also enables other offloading features such as TCP segmentation offloading or IPSec offloading. Finally, even if the CRC/checksum computation is fixed and not programmable in current NICs, it seems not necessary to develop more flexible and programmable engines for checksum computation due to the very limited interest in deploying checksum algorithms different from those used today.

2) Encryption/Decryption: The encryption/decryption of packet payloads requires very expensive computations, and as such, it is a good candidate for NIC offloading. Due to the large number of encryption standards, each one with its specific algorithm and implementation characteristics, it is nearly impossible[2] to have a generic or programmable encryption offloading engine. Instead, there are several NICs that are able to provide specific protocol offloading, such as IPSec, with a subset of encryption algorithms [e.g., Advanced Encryption Standard—Cipher Block Chaining (AES-CBC)] with authentication features (e.g., HMAC-SHA1 and HMAC-SHA2 with a different key length), or more recent Authenticated Encryption with

---

[2]Nor advisable! As is well known, implementation of cryptographic primitives requires professional developers with extremely specific skills, e.g., to avoid secret information leakage via side channels that a general-purpose developer may inadvertently introduce in the design.

Associated Data algorithms such as AES-GCM or AES-CCM. In some cases, the encryption offloading can save more than 50% of CPU time, providing significant gains of the packet forwarding throughput [82]. However, due to the complexity of the hardware development of encryption algorithms, only high-end NICs provide IPSec offloading features. Moreover, integrating the encryption engine in the ASIC chip providing the basic NIC functionalities can be very costly. A possible solution utilized by several vendors is to use a network processing unit (NPU) or an FPGA inside the NIC as the device in which the encryption engine is realized. This choice presents several benefits, since it is possible to add novel algorithms updating the NIC firmware, and it makes very easy to provide patches if security problems arise in the implementation of the specific protocol. From the programmability point of view, this is an opportunity for the development of NFV, since the presence of a programmable device in the NIC could permit a more timely network security upgrades and a much faster patching of emerging vulnerabilities. In short, the hardware designed to provide programmable encryption engine already represents an embryo of an FPGA enabled SmartNIC, as those described next.

3) Large Send/Receive Offloading: Transport protocols usually split the stream of data to transmit into smaller chunks with a maximum size given by the network infrastructure. The task of splitting a stream in smaller segments, as well as the receiver task of rebuilding the stream from the received packet, can consume several CPU clocks. In particular, there is a considerable overhead caused by the system call associated with each packet. These overheads can be avoided if the task of splitting/rebuilding the packets is offloaded to the NIC. In case of data sending, this is generically called generic segmentation offload and it is called TCP segmentation offload when applied to TCP. The offloading engine breaks the data coming from the host into smaller segments, adds the TCP, IP, and data link layer protocol headers to each segment, recomputes the checksum, and sends the packet over the network. For the receive side, the offloading is called large receive offload (LRO) and allows the system to reduce the CPU time for processing TCP packets that arrive from the network and to reduce the interrupt pressure from the NIC. We notice that these techniques complement the I/O batching technique early seen; whereas I/O batching maintains the individual packets received by the card unaltered and reduces interrupt pressure, segmentation offloading goes one step further by offering to disassemble/reassemble higher layer data into/from several segments, which reduces interrupt pressure. Furthermore, LRO can also significantly increase the throughput of software network functions. In fact, in software NFs, most of the processing overhead comes from per-packet operations. Increasing the packet size well above the maximum transmission unit (MTU) size (a reassembled packet could have up to packet 64 kB/s) drastically reduces the overhead, thereby increasing the overall throughput of a software network function.

Still, it is worth to remark that when, as in the case of TCP, packet batching occurs at the NIC level, the subsequent transmission of batched packets may occur in (micro) bursts, which may increase packet loss in some scenarios [83]. Programmable pacing techniques directly implemented in the NIC HW are an interesting research direction to address such issue [84].

Another hot topic for segment offloading is related to the use of this technique when UDP is used to build an overlay protocol, e.g., for QUIC. Since each QUIC packet in its encrypted header a packet number (PN) identifier, the segmentation/reassembly of UDP packets breaks the correspondence between the UDP packet and the PN. Due to the encryption, it is not straightforward to update the PN during the segmentation/reassembly operations. Thus, native QUIC segmentation offloading can be done combining the HW encryption with a modified HW for segmentation offloading. However, as will be discussed extensively in Section V-C, designing such offloading hardware depends on how largely adopted will be this protocol.

B. SmartNIC Architectures

A SmartNIC hosts a programmable engine that allows for the implementation and extension of the NIC's functions, e.g., for hardware offloading of new protocols that become of widespread adoption. Almost all the SmartNICs work as a bump-in-the-wire device inserted between the PHY Ethernet interface and the host Peripheral Component Interconnect Express (PCIe) interface [85]. This type of SmartNIC usually provides several network interfaces, a PCIe interface to communicate with the host, and different memory peripherals such as external DRAMs and SRAMs. The former are bigger memories with considerable latency and are used for packet buffering, and the latter is lower latency and is used to store random access information (e.g., hash or routing tables, counters, and so on). The core of the SmartNIC is a programmable element that can be a multicore system-on-chip (SoC) or an FPGA, as discussed next.

1) Multicore SoC-Based SmartNICs: Examples of multicore SoCs are the Cavium [86] and BlueField [87], based on general-purpose CPU cores (mainly ARM cores), or the Netronome [88] and Tilera [89], based on smaller cores specifically designed for networking applications. Usually, these SoCs also have several hardware blocks, usually called acceleration engines (AEs) that perform specific network tasks such as encryption or packet classification (using on-chip TCAMs). These processor-based SmartNICs also use software acceleration techniques such as those described in Section III (some of them can run applications on top of the DPDK framework) but at the same time share the same drawbacks in terms of latency and latency variability as well as memory contention, state synchronization, and load balancing issues among cores. In summary, these devices allow for a good degree of pro-

grammability but offer limited scalability and maximum achievable throughput, e.g., it is usually hard to scale their performance beyond 50–100 GbE [90].

2) FPGA-Based SmartNICs: The other solution for the realization of a SmartNIC is an FPGA-based card. A well-known example of this kind of SmartNIC, in the academic domain, is the NetFPGA SUME [67], which provides 4×10 GbE SFP+ links, a x8 Gen3 PCIe interface, a Xilinx Virtex-7 XC7V690T FPGA, three 72-Mb SRAM chips, and two 4-GB DDR3 DRAMs. Several other cards with similar characteristics are available in the market. These cards differ for the FPGA performance and size and for the type of attached peripheral (number and speed of network links, size of external memories, and so on). The performance of FPGA-based SmartNICs is extremely competitive in terms of latency, power consumption, and cost as extensively discussed in [90].

C. Programming Abstractions

SmartNICs differ not only in their hardware architecture but also on the way that the hardware features can be accessed, programmed, and extended. In the most general case, an FPGA-based SmartNIC can be programmed with Hardware Description Languages (such as VHDL and Verilog), whereas Multi-SoC-based SmartNICs offer programmability through some high-level language (e.g., a C dialect used to program specialized NPs). However, since these languages require a programmer to have extensive hardware knowledge (either for FPGA or NPs), domain-specific abstractions for network functions are often also provided. We can categorize these abstractions in stateless and stateful ones. A stateful abstraction allows a programmer to describe algorithms that read and modify state, whereas a stateless abstraction can only read state that has been provided to the device through some other means, e.g., written by a nonprogrammable block.

1) Stateless Abstractions: The match-action table (MAT) abstraction is widely used to describe a stateless packet processing data plane. The abstraction allows a programmer to define which actions should be applied to the packets belonging to a given network flow. The flow is specified by the match part of MAT's entry, while the corresponding actions can forward or drop a packet and modify its header. The entries of the MAT cannot be changed using the abstraction itself. It is assumed that a separated control plane asynchronously updates the entries. In other words, the entries' update operations are not executed in the device's fast path.

MATs are used in several contexts. OpenFlow [91] adopts a pipeline of MATs to describe packet processing in network switches. Some SmartNICs implement OpenFlow-like abstractions, and even regular NICs may offer a limited MAT abstraction support [92]. Reconfigurable match tables (RMTs) [93] extend the OpenFlow approach with the ability to define the type and number of MATs, including the description of custom packet headers and actions.

Originally targeted to hardware switches (P4 switches may indeed leverage an internal architecture based on the RMT technology), the approach has also been recently applied to SmartNICs [69], [74], [75]. Given their stateless nature, MAT-based abstractions are generally used only to describe a small subset of functional blocks, such as header read and write and a small set of flow actions.

2) Stateful Abstractions: While MAT abstractions dominate the scene for stateless packet processing, it is still unclear which abstraction is most appropriate for capturing high-performance stateful processing. Often, a very generic approach is used; the programmer writes tiny packet programs that are executed for each processed packet. The programs are typically expressed with a restricted programming language in order to reduce the risk of negatively impacting performance (e.g., removing unbounded loops) and to increase system safety (e.g., forbidding flexible, pointer-based memory accesses).

This is the approach used by packetC [94], protocol-oblivious forwarding (POF) [95], Linux's eBPF [96] and Domino [97], which mainly differ for the target platforms and for the state access model. For instance, packetC and POF target NPs, and eBPF is usually executed within the Linux kernel but can be offloaded to some SmartNICs, while Domino targets high-performance ASIC implementations. In terms of state management, these abstractions are influenced by the target executor. Thus, when the target is a multiprocessor system (NPs and multicore CPUs), state is partitioned in processor-local and global parts, and the programmer has to carefully manage state access. In the case of Domino, the state is always global, and state consistency and performance are guaranteed at compile time at the cost of limited flexibility.

A different, more structured, abstraction was originally suggested by OpenState [98] and FAST [99]. In both cases, packets are grouped in programmer-defined network flows (such as in MATs), and flows are associated with a finite-state machine (FSM). Programmers define the FSMs transition tables, which are used to update the corresponding FSMs' states as packets are processed. When processing a packet, the corresponding FSM's current state is used to select the packet forwarding and modification actions. The compeling nature of such abstractions is their very efficient implementation in hardware, in a small and deterministic number of clock cycles—a state transition, in fact, can be modeled and executed via a suitably extended flow table [98]. Moreover, OpenState was extended in [100] with the support for triggering conditions, flow context memory registers, and on-the-fly arithmetic/logic register update functions. These stateful abstractions allow a programmer to describe more complex functions, thus supporting also the offloading of more demanding flow state tracking and, in some cases, session management blocks.

At last, please note that all the stateful HW programming abstractions mentioned earlier have been designed

for the somewhat "simpler" context of programmable data planes inside network switches, i.e., on-the-fly forwarding and packet-level processing functions. Even in such scenario, to date, it is still unclear what is the "right" compromise between flexibility, performance, and openness. To a greater extent, such compromise becomes critical when we seek programming abstraction focused on higher layer network functions, which may involve network functions well beyond packet-level "on-the-fly" operations—hence further accounting for protocol states, timers, opportunistic packet buffering, and signaling messages. The quest for a sufficiently expressive programming abstraction, tailored to permit offloading in hardware NICs of transport/application layer functions or even implement way more challenging cloudlike functions such as those posited in [101], is still open and appears to be an exciting research topic.

V. KEY DERIVATIONS AND GUIDELINES

In this section, we present the guidelines concerning the different techniques that were detailed in Sections II–IV. They include general guidelines and caveats that should be taken into account when designing and developing VNFs (Section V-A). Furthermore, we provide the guidelines that are specific to the two main categories of acceleration techniques that are covered in this paper, i.e., software (Section V-B) and hardware (Section V-C) acceleration. Finally, we demonstrate how these guidelines can be applied to the exemplary NGFW function (Section V-D).

A. Guidelines for Design and Development of Network Functions

Ideally, as outlined in Section II, the VNF ecosystem should be designed based on a detailed performance requirement specification for the intended deployment scenario. However, due to internal policies or previous design choices, the essential execution environment is often already given when the performance becomes constraining. In the following, we concentrate on this deployment optimization scenario when the task is to improve VNF performance for a given environment. We highlight the best practices for the VNF adoption, derive the corresponding suggested implementation patterns, and summarize caveats in antipatterns, which should be avoided.

1) Know Relevant Performance Metrics: To bootstrap the performance optimization process, it is necessary to be aware of the network functions, the current workload, the resource utilization of the whole ecosystem, as well as the underlying network topology.

✓ *Pattern:* Deployment of dedicated monitoring and verification tools providing up-to-date reports about the actual performance and resource descriptors in the system.

✗ *Antipattern:* Inaccurate view due to missing monitoring information and insufficient insights may result in resource overprovisioning, e.g., by vertical scaling.

2) Optimize for the Most Constraining Resource: Apply the "Pareto principle" of software optimization to identify the 10% of the code affecting the most critical resource (e.g., most of the running time is spent and most of the memory is used) and apply the previously outlined software and hardware acceleration techniques. Table 1 is a good starting point to identify the bottleneck resources for the specific VNF type.

✓ *Pattern:* Repeat the following performance optimization loop.

1) Identify per-VNF critical resources.
2) Obtain deployment-specific insights using monitoring tools.
3) Optimize for the most constraining resource by choosing appropriate software or hardware acceleration.
4) Test and evaluate the results.

✗ *Antipattern:* Rewriting a VNF for using a specific acceleration technique will improve one specific bottleneck resource. For instance, [102] shows that SGW acceleration with DPDK improves raw packet-per-second performance, while user-space table lookup remains a bottleneck, resulting in poor performance as soon as the table is sufficiently populated.

3) Exploit Parallelism: Network-related workloads typically lend themselves readily to a parallel pipelined execution model, by providing ample opportunities to separate a workload to multiple threads at the packet level, the protocol level, or the network function level. Network-related workloads can be balanced among multiple threads at the packet, protocol, or network function level. Thus, it is possible to utilize the parallel execution capabilities of modern hardware devices to execute atomic building blocks in parallel.

✓ *Pattern:* Build and optimize network functions to maximize concurrence by minimizing contention on shared resources and point-to-point synchronization (coherency) between execution threads.

✗ *Antipattern:* Writing code that is fundamentally serial in nature causes significant runtime overhead due to synchronization between execution threads. Similarly, resource-hungry synchronization primitives or frequent memory accesses across NUMA nodes constitute performance bottlenecks.

4) Use Lightweight Virtualization and Isolation: The choice of the specific virtualization technique is a careful balancing act between flexibility, performance, overhead, and degree of isolation with respect to CPU, network, and storage resources.

✓ *Pattern:* Use of the lightest possible isolation technique that is still sufficient to provide the required isolation barrier. This might also include higher layer isolation techniques for safe zero-cost isolation. Examples are compile-time-type checking and a safe runtime enabled purely in software [103].

✗ *Antipattern:* Using lightweight virtualization across trust boundaries, e.g., deploying sensitive applications of different tenants into containers with limited isolation capabilities.

5) Cautious Use of Hardware Offloading: Hardware offloading is very appealing since it allows to perform certain CPU-intensive operations in hardware, thereby relieving the CPU. Possible drawbacks, however, are a lack of isolation as well as additional I/O costs. Furthermore, the network topology should be taken into account since hardware offloading capabilities might be limited to a subset of network nodes.

✓ *Pattern:* Careful measurement of critical and reusable software building blocks that can be effectively realized in hardware and a step-by-step process to apply offloading gradually to CPU-intensive code.

✗ *Antipattern:* Using hardware offloading in the middle of a service chain may involve an unnecessary I/O roundtrip to hardware, possibly slowing execution down instead of accelerating it.

B. Guidelines for Software Acceleration

Based on our expertise in the software networking domain, we provide the guidelines for the development of VNFs. From the point of view of a developer, two points should be kept in mind. First, a tradeoff may arise between the acceleration benefit and the complexity of the implementation of a specific acceleration technique. Second, some of the acceleration techniques work better in conjunction with some classes of VNFs. We provide a brief list of good practices and develop each point in the following.

1) Utilize Hardware-Supported Functions When Available: Hardware-supported functions can have a significant performance impact and can be easily introduced into the existing software due to widespread driver integration. The two most important examples from this category include hardware counters and RSS hash functions. While using the former can significantly reduce the amount of state and memory accesses, the latter lowers the number of CPU cycles.

2) Build Virtual Network Functions on Top of General-Purpose-Accelerated Frameworks: Nowadays, high-speed frameworks exploit several state-of-the-art acceleration techniques and also provide mechanisms for interconnecting several VNFs at high speed. Hence, developers can focus on their particular application while leveraging the infrastructure and flexibility that is offered by such frameworks. The choice of the specific framework in the wide ecosystem may be dictated by the environment (e.g., preexisting expertise), by specific framework capabilities (e.g., availability of specific functions), or finally by performance considerations (i.e., which is harder as of now, since a comprehensive comparison of the available frameworks is still an open research question).

3) Implement Very Specific Tasks Using Low-Level Building Blocks for CPU-Intensive Tasks and Leverage Advanced Coding Styles: Finally, when computation-intensive VNFs have to be deployed on bare metal or on top of an existing high-speed framework, there may be cases in which the I/O is dominated by the computation or situations where the bottleneck of I/O is negligible. For an overloaded system, i.e., I/O and CPU bottlenecks in cascade, polling and I/O batching can accelerate the packet processing. More generally, memory-related acceleration, compute batching, and high-level programming techniques are beneficial for the VNF (paying attention to the zero-copy approach, which can be useful when there is no further payload modification). Collectively, the use of these techniques allows to optimize CPU processing of VNFs in software.

C. Guidelines for Hardware Acceleration

Due to the required amount of expert knowledge in different domains such as programming languages, hardware architectures, and verification mechanisms, the design of hardware-accelerated functions requires significantly more effort than the corresponding process for their software counterpart. Therefore, the choice of offloading a network function, or part of it, to a SmartNIC is a careful decision that depends on several factors that we summarize in the following.

1) Performance Gain: Offloading a primitive into the NIC is recommended when the processing power, e.g., in terms of CPU cycles, saved by doing so is significant.

2) Primitive Adoption: The widespread use of a primitive is an important factor when deciding whether it is worth to implement it in the NIC.

3) Complexity: Due to the high cost of nonrecurring engineering of hardware development, the decision about the offloading heavily depends on the complexity of the primitive. Furthermore, the expected frequency of updates to the functionality of network functions should be taken into account due to the fact that each update induces a development overhead. Finally, the presence of temporal dynamics that require frequent migration of VNFs should be carefully considered since the migration process tends to be more complex with hardware-accelerated components in contrast to purely software-based ones.

4) Relation of the Primitive to Other Processing Functions: Offloading to the NIC is only useful if the subsequent operations are also offloaded to the NIC. Otherwise, costly back and forth movements between the NIC and the host are required.

Finally, it is worth noting that the use of network-focused programming abstractions for the design of hardware-accelerated functions can greatly simplify the development cycle. In our opinion, this is an important and still immature research field that could provide a significant advance in network development in the future. This is particularly true when dealing with issues associated with

running multiple applications using the same accelerator and multitenancy.

SmartNICs are one of the most discussed and evolving technologies for accelerating network functions. As mentioned in Section IV-B, there are two main SmartNIC types: multicore SoC-based and FPGA-based.

Designing efficient network functions using Multicore NICs requires to deal with some typical issues of parallel programming. The single core of the SoC can be programmed using standard programming languages (mainly C/C++ is used) and does not present specific challenging issues. On the other hand, the interaction among the cores, the access to shared resources, and the efficient use of the AEs as well as the distribution of the packet among the cores are hard to manage and often require a deep knowledge of the specific hardware architecture of the SoC. This makes the implementation very target specific, preventing the porting to a different SoC-based SmartNIC and slows down the development of efficient network functions.

When FPGA-based SmartNICs are programmed using low-level HDL, they demand extremely long development and verification time. Thus, even if, in terms of absolute performance, they could provide better throughput than SoC-based SmartNICs, the design effort could make it more convenient to use more programmable friendly SoC-based SmartNICs. If the programming abstractions described in Section IV-C will evolve to mature and easy to use programming paradigm, the FPGA-based SmartNIC will emerge as the right element to provide programmable hardware accelerators for NFs.

From the performance point of view, SmartNICs, in particular FPGA-based ones, provide benefits for both latency and throughput. For instance, [90] reports a 3–5× improvement in terms of average and worst case latency as well as a sixfold increase in throughput when employing an FPGA-based SmartNIC. While these gains have to be interpreted as of anecdotal value, since they relate about a performance comparison of a specific design, they nevertheless are useful to report, to perceive the expected gains in "order sense" than the offloading of complex functions to a SmartNIC can bring. In terms of economic advantage in using SmartNICs, one should consider that such devices can free valuable CPU cores for other tasks. For instance, Microsoft reports that, in a cloud datacenter, each core has an equivalent value of about $800 per year [90].

Finally, there are cases in which using an FPGA-based SmartNIC provides much larger benefits. This is the case, for instance, of applications that require very low-latency and/or predictable performance, such as high-performance trading. Another area where SmartNICs can be beneficial is deployments that are somehow space constrained, e.g., in an edge datacenter deployed in a big city, where real estate costs are particularly expensive. In these cases, either the other approaches are completely ruled out by unmet requirements (e.g., latency below 1us) or the application-level and the overall infrastructure-level advantages offset the extra cost paid in buying a powerful SmartNIC.

D. Potential of Acceleration Techniques for the NGFW

In addition to the above-mentioned guidelines regarding software- and hardware-based acceleration techniques, we present an overview of insights that are relevant to the NGFW VNF. To this end, we first recall the building blocks that are associated with this VNF in Table 1 and then illustrate how we can improve their performance. We note, however, that since the main focus of this paper is not the NGFW, we do not provide an exhaustive set of guidelines and options but exemplarily demonstrate how three building blocks might be improved.

On the one hand, the NGFW comes with cryptography support and therefore needs to perform compute-intensive decryption and encryption tasks on packet payloads. As outlined in Section IV-A2, significant amounts of CPU time can be saved by offloading these tasks to a hardware entity. Depending on the capabilities in terms of supported encryption standards, this entity might be a commodity NIC, a specialized NPU, or an FPGA component inside the NIC.

On the other hand, the NGFW has NAPT and VPN capabilities that require rewriting packet headers. In this context, switching from plain virtualization to software acceleration via frameworks, such as DPDK, has been shown to significantly improve performance [104]. In the referenced work, we compare the performance of a long-term evolution (LTE) SGW VNF that is implemented as a user-space application with a DPDK-accelerated version of the same VNF. We demonstrate a throughput improvement of almost factor 8 when using the latter. These results can be translated to the NGFW since both the functions share the corresponding building block.

Furthermore, the NGFW employs DPI techniques in its decision-making process which require tracking flow and session states as well as processing header and payload information. Whereas, as previously introduced, we deem a thorough account of DPI acceleration via FPGA or GPUs to be out of the scope of this paper, it is nevertheless instructive to consider how the software guidelines expressed earlier can boost the performance of a "vanilla" DPI application in the NGFW. Indeed, even though no packet modification is performed during these tasks, significant improvements in terms of maximum throughput are possible when using SR-IOV-based virtualization in conjunction with an acceleration framework such as DPDK [102]. When compared with an implementation based on the default Linux kernel stack and a libpcap-based DPI application, an up to tenfold throughput increase is observed. While, again, these gains have to be interpreted as of anecdotal value, they relate the expected gains in "order sense" than adopting state-of-the-art software techniques can bring.

VI. CONCLUSION

The rising amount of networked devices and services, the heterogeneity of their demands and communication patters, as well as their temporal dynamics pose stringent requirements that exceed the capabilities of today's networking architectures. These phenomena call for a new type of scalable and adaptive systems. Network softwarization, which comprises networking paradigms such as SDN and NFV, is seen as a possible solution to fulfill these requirements and to provide a scalable and adaptive networking infrastructure. As NFV provides the necessary degree of flexibility by moving today's embedded network functions to software running on common server hardware, recent efforts try to improve the performance of these software-based solutions.

In this paper, we conduct a survey of the evolving research field of performance acceleration techniques for NFV. First, we categorize a subset of representative network functions with respect to common low-level building blocks. In conjunction with information on the resource usage of individual building blocks, we identify the corresponding performance bottlenecks.

Then, we review and summarize software-based and hardware-based acceleration techniques in a holistic manner. We first highlight pure software acceleration techniques allowing to speed up the performance of a network function itself. Then, we cover the software stack as a whole and discussed the interplay between software and driver-based hardware access. In the case of hardware acceleration techniques, we cover offloading strategies based on standard NICs, programmable SmartNICs, as well as domain-specific programming abstractions. These mechanisms represent tradeoffs with respect to the implementation overhead, performance gains, and costs.

Based on the survey, we finally derive guidelines that allow the reader to understand which acceleration techniques might help to improve the performance of specific network functions. Our summary includes generic guidelines that represent best practices when optimizing the performance of a given deployment scenario, as well as specific guidelines on how to apply software and hardware acceleration techniques. This paper thus gives guidance to potential NFV adopters on how to improve the performance of their current system, clearly highlighting the research direction that is still open and requires a community wide attention. ∎

REFERENCES

[1] "Network functions virtualisation (NFV); use cases," Tech. Rep. ETSI GR NFV 001, Rev. 1.2.1, 2017.

[2] B. Yi, X. Wang, K. Li, S. K. Das, and M. Huang, "A comprehensive survey of network function virtualization," *Comput. Netw.*, vol. 133, pp. 212–262, Mar. 2018.

[3] B. Carpenter and S. Brim, *Middleboxes: Taxonomy and Issues*, document RFC 3234, 2002.

[4] M. Hoffmann *et al.*, "SDN and NFV as enabler for the distributed network cloud," *Mobile Netw. Appl.*, vol. 23, no. 3, pp. 521–528, 2018.

[5] *Network Functions Virtualisation (NFV); NFV Performance & Portability Best Practises*, document ETSI GS NFV-PER 001 V1.1.1 Rev. 1.1.2, 2014.

[6] *The Bro Network Security Monitor*. [Online]. Available: https://www.bro.org/

[7] *Snort IDS*. [Online]. Available: https://www.snort.org/

[8] M. Trevisan, A. Finamore, M. Mellia, M. Munafo, and D. Rossi, "Traffic analysis with off-the-shelf hardware: challenges and lessons learned," *IEEE Commun. Mag.*, vol. 55, no. 3, pp. 163–169, Mar. 2017.

[9] V. Addanki, L. Linguaglossa, J. Roberts, and D. Rossi, "Controlling software router resource sharing by fair packet dropping," *IFIP Netw.*, 2018.

[10] L. Deri, M. Martinelli, T. Bujlow, and A. Cardigliano, "nDPI: Open-source high-speed deep packet inspection," in *Proc. Int. Wireless Commun. Mobile Comput. Conf. (IWCMC)*, Aug. 2014, pp. 617–622.

[11] *Network Functions Virtualisation (NFV); Acceleration Technologies*, document ETSI GS NFV-IFA 001 V1.1.1, 2015.

[12] J. Sahoo, S. Mohapatra, and R. Lath, "Virtualization: A survey on concepts, taxonomy and associated security issues," in *Proc. Int. Conf. Comput. Netw. Technol. (ICCNT)*, Apr. 2010, pp. 222–226.

[13] P. Barham *et al.*, "Xen and the art of virtualization," *ACM SIGOPS Operating Syst. Rev.*, vol. 37, no. 5, pp. 164–177, 2003.

[14] *VMWare ESXi*. [Online]. Available: https://www.vmware.com/products/esxi-and-esx.html

[15] A. Kivity, Y. Kamay, D. Laor, U. Lublin, and A. Liguori, "KVM: The Linux virtual machine monitor," in *Proc. Linux Symp.*, 2007, pp. 225–230.

[16] J. Hwang, S. Zeng, F. Y. Wu, and T. Wood, "A component-based performance comparison of four hypervisors," in *Proc. IM*, May 2013, pp. 269–276.

[17] G. Lettieri, V. Maffione, and L. Rizzo, "A survey of fast packet I/O technologies for network function virtualization," in *Proc. Int. Conf. High Perform. Comput.*, 2017, pp. 579–590.

[18] R. Russell, "Virtio: Towards a de-facto standard for virtual I/O devices," *ACM SIGOPS Operating Syst. Rev.*, vol. 42, no. 5, pp. 95–103, 2008.

[19] V. Maffione, L. Rizzo, and G. Lettieri, "Flexible virtual machine networking using netmap passthrough," in *Proc. IEEE Int. Symp. Local Metropolitan Area Netw. (LANMAN)*, Jun. 2016, pp. 1–6.

[20] D. Merkel, "Docker: Lightweight Linux containers for consistent development and deployment," *Linux J.*, vol. 239, p. 4, Mar. 2014.

[21] P.-H. Kamp and R. N. Watson, "Jails: Confining the omnipotent root," in *Proc. Int. SANE Conf.*, 2000, p. 116.

[22] A. Madhavapeddy and D. J. Scott, "Unikernels: The rise of the virtual library operating system," *Commun. ACM*, vol. 57, no. 1, pp. 61–69, 2014.

[23] G. McGrath and P. R. Brenner, "Serverless computing: Design, implementation, and performance," in *Proc. Int. Conf. Distrib. Comput. Syst. Workshops (ICDCSW)*, Jun. 2017, pp. 405–410.

[24] I. Baldini *et al.*, "Serverless computing: Current trends and open problems," in *Research Advances in Cloud Computing*. 2017.

[25] N. M. K. Chowdhury and R. Boutaba, "A survey of network virtualization," *Comput. Netw.*, vol. 54, pp. 862–876, Apr. 2010.

[26] A. Blenk, A. Basta, M. Reisslein, and W. Kellerer, "Survey on network virtualization hypervisors for software defined networking," *IEEE Commun. Surveys Tuts.*, vol. 18, no. 1, pp. 655–685, 1st Quart., 2016.

[27] *Network Functions Virtualisation (NFV); Architectural Framework*, document ETSI GS NFV 002 V1.1.1 (2013-10), 2014.

[28] N. F. S. de Sousa, D. A. L. Perez, R. V. Rosa, M. A. Santos, and C. E. Rothenberg. (2018). "Network service orchestration: A survey." [Online]. Available: https://arxiv.org/abs/1803.06596

[29] H. Guan, Y. Dong, R. Ma, D. Xu, Y. Zhang, and J. Li, "Performance enhancement for network I/O virtualization with efficient interrupt coalescing and virtual receive-side scaling," *IEEE Trans. Parallel Distrib. Syst.*, vol. 24, no. 6, pp. 1118–1128, Jun. 2013.

[30] L. Rizzo, "Device polling support for FreeBSD," in *Proc. BSDConEurope Conf.*, 2001.

[31] *The Linux Foundation*. [Online]. Available: https://wiki.linuxfoundation.org/networking/napi

[32] S. Han, K. Jang, K. Park, and S. Moon, "PacketShader: A GPU-accelerated software router," in *Proc. SIGCOMM*, 2010, pp. 195–206.

[33] J. Kim, S. Huh, K. Jang, K. Park, and S. Moon, "The power of batching in the click modular router," in *Proc. Asia–Pacific Workshop Syst.*, 2012, p. 14.

[34] L. Linguaglossa *et al.*, "High-speed software data plane via vectorized packet processing (extended version)," Telecom ParisTech, Tech. Rep., 2017.

[35] *Data Plane Development Kit*. [Online]. Available: http://dpdk.org

[36] L. Rizzo, "netmap: A novel framework for fast packet I/O," in *Proc. USENIX ATC*, 2012.

[37] L. Deri. [Online]. Available: https://www.ntop.org/products/packet-capture/pf_ring/pf_ring-zc-zero-co%py/

[38] G. Rétvçri, J. Tapolcai, A. Kőrösi, A. Majdán, and Z. Heszberger, "Compressing IP forwarding tables: Towards entropy bounds and beyond," *ACM SIGCOMM Computer Commun. Rev.*, vol. 43, no. 4, pp. 111–122, 2013.

[39] T. Barbette, C. Soldani, and L. Mathy, "Fast userspace packet processing," in *Proc. ANCS*, May 2015, pp. 5-16.

[40] L. Molnár *et al.*, "Dataplane specialization for high-performance openflow software switching," in *Proc. ACM SIGCOMM*, 2016, pp. 539–552.

[41] A. Beaupré. (2017). *New Approaches to Network Fast Paths*. [Online]. Available: https://lwn.net/Articles/719850/

[42] *eXpress Data Path (XDP)*. [Online]. Available: https://www.iovisor.org/technology/xdp

[43] J. Corbet, "Zero-copy networking," Linux Weekly News, Tech. Rep., 2017.

[44] A. Kalia, D. Zhou, M. Kaminsky, and D. G. Andersen, "Raising the bar for using GPUs in software packet processing," in *Proc. NSDI*, 2015.

[45] R. Mansilha *et al.*, "Hierarchical content stores in high-speed ICN routers: Emulation and prototype implementation," in *Proc. ICN*, 2015, pp. 59–68.

[46] G. Xylomenos *et al.*, "A survey of information-centric networking research," *IEEE Commun. Surveys Tuts.*, vol. 16, no. 2, pp. 1024–1049, 2nd Quart., 2014.

[47] I. Marinos, R. N. M. Watson, and M. Handley, "Network stack specialization for performance," in *Proc. SIGCOMM*, 2014, pp. 175–186.

[48] T. Zhang, L. Linguaglossa, M. Gallo, P. Giaccone, and D. Rossi, "FlowMon-DPDK: Parsimonious per-flow software monitoring at line rate," in *Proc. IFIP Traffic Monit. Anal.*, 2018.

[49] E. Jeong *et al.*, "mTCP: A highly scalable user-level TCP stack for multicore systems," in *Proc. NSDI*, 2014.

[50] B. Pfaff *et al.*, "The design and implementation of Open vSwitch," in *USENIX NSDI*, 2015.

[51] E. Kohler, R. Morris, B. Chen, J. Jannotti, and F. Kaashoek, "The Click Modular Router," *Operating Syst. Rev.*, vol. 18, no. 3, pp. 263–297, 1999.

[52] S. Han, K. Jang, A. Panda, S. Palkar, D. Han, and S. Ratnasamy, "SoftNIC: A software NIC to augment hardware," EECS Department, Univ. California, Berkeley, CA, USA, Tech. Rep. UCB/EECS-2015-155, 2015.

[53] J. Martins *et al.*, "ClickOS and the art of network function virtualization," in *Proc. USENIX Conf. Netw. Syst. Design Implement.*, 2014.

[54] M. Gallo and R. Laufer, "ClickNF: A modular stack for custom network functions," in *Proc. USENIX Annu. Tech. Conf. ATC*, 2018.

[55] M. A. Jamshed, Y. G. Moon, D. Kim, D. Han, and K. Park, "mOS: A reusable networking stack for flow monitoring middleboxes," in *Proc. NSDI*, 2017.

[56] A. Abdelrazik, G. Bunce, K. Cacciatore, K. Hui, S. Mahankali, and F. Van Rooyen, "Adding speed and agility to virtualized infrastructure with openstack," Tech. Rep., 2017.

[57] J. Hwang, K. K. Ramakrishnan, and T. Wood, "NetVM: High performance and flexible networking using virtualization on commodity platforms," *IEEE Trans. Netw. Service Manage.*, vol. 12, no. 1, pp. 34–47, Mar. 2015.

[58] A. Belay *et al.*, "The IX operating system: Combining low latency, high throughput, and efficiency in a protected dataplane," *ACM Trans. Comput. Syst.*, vol. 34, no. 4, p. 11, 2017.

[59] S. Peter, T. Anderson, and T. Roscoe, "Arrakis: The operating system as control plane," *ACM Trans. Comput. Syst.*, 2013.

[60] N. Pitaev, M. Falkner, A. Leivadeas, and I. Lambadaris, "Characterizing the performance of concurrent virtualized network functions with OVS-DPDK, FD. IO VPP and SR-IOV," in *Proc. ACM/SPEC Int. Conf. Perform. Eng.*, 2018, pp. 285–292.

[61] N. Zilberman, P. M. Watts, C. Rotsos, and A. W. Moore, "Reconfigurable network systems and software-defined networking," *Proc. IEEE*, vol. 103, no. 7, pp. 1102–1124, Jul. 2015.

[62] Intel. *Intel 10 Gigabit AT Server Adapter*. [Online]. Available: https://www.intel.com/content/dam/support/us/en/documents/network/adapt%er/10gbe/atserver/sb/318349.pdf

[63] Mellanox. *ConnectX-6 EN IC 200Gb/s Ethernet Adapter IC*. [Online]. Available: http://www.mellanox.com/related-docs/prod_silicon/PB_ConnectX-6_EN_IC.p%df

[64] *Network Functions Virtualisation (NFV);Acceleration Technologies*, document ETSI GS NFV-IFA 002 V2.1.1, 2017.

[65] *openNFP*. (2016). [Online]. Available: http://open-nfp.org/

[66] Netronome. *AgilioTM CX 2x40GbE Intelligent Server Adapter*. [Online]. Available: https://www.netronome.com/media/redactor_files/PB_Agilio_CX_2x40GbE.pdf

[67] N. Zilberman, Y. Audzevich, G. A. Covington, and A. W. Moore, "NetFPGA SUME: Toward 100 Gbps as research commodity," *IEEE Micro*, vol. 34, no. 5, pp. 32–41, Sep./Oct. 2014.

[68] L. Jose, L. Yan, G. Varghese, and N. McKeown, "Compiling packet programs to reconfigurable switches," in *Proc. USENIX Symp. Netw. Syst. Design Implement. (NSDI)*, 2015.

[69] A. Kaufmann, S. Peter, T. Anderson, and A. Krishnamurthy, "FlexNIC: Rethinking network DMA," in *Proc. USENIX HotOS*, 2015.

[70] S. Li *et al.*, "Architecting to achieve a billion requests per second throughput on a single key-value store server platform," in *Proc. Annu. 42nd Int. Symp. Comput. Archit.*, Jun. 2015, pp. 476–488.

[71] X. Li, R. Sethi, M. Kaminsky, D. G. Andersen, and M. J. Freedman, "Be fast, cheap and in control with switchkv," in *Proc. USENIX Symp. Netw. Syst. Design Implement. (NSDI)*, 2016.

[72] H. Ballani *et al.*, "Shea, "Enabling end-host network functions," *SIGCOMM Comput. Commun. Rev.*, vol. 45, no. 4, pp. 493–507, 2015.

[73] D. Firestone, "VFP: A virtual switch platform for host SDN in the public cloud," in *Proc. USENIX Symp. Netw. Syst. Design Implement. (NSDI)*, 2017, pp. 315–328.

[74] H. Wang *et al.*, "P4FPGA: A rapid prototyping framework for p4," in *Proc. Symp. SDN Res.*, 2017, pp. 122–135.

[75] G. Brebner and W. Jiang, "High-speed packet processing using reconfigurable computing," *IEEE Micro*, vol. 34, no. 1, pp. 8–18, Jan./Feb. 2014.

[76] G. Vasiliadis, M. Polychronakis, S. Antonatos, E. P. Markatos, and S. Ioannidis, "Regular expression matching on graphics hardware for intrusion detection," in *Proc. Int. Workshop Recent Adv. Intrusion Detection*, 2009, pp. 265–283.

[77] S. Pontarelli, G. Bianchi, and S. Teofili, "Traffic-aware design of a high-speed FPGA network intrusion detection system," *IEEE Trans. Comput.*, vol. 62, no. 11, pp. 2322–2334, Nov. 2013.

[78] S. Dharmapurikar, P. Krishnamurthy, T. S. Sproull, and J. W. Lockwood, "Deep packet inspection using parallel bloom filters," in *Proc. Symp. High Perform. Interconnects*, Aug. 2003, pp. 44–51.

[79] F. Yu, Z. Chen, Y. Diao, T. V. Lakshman, and R. H. Katz, "Fast and memory-efficient regular expression matching for deep packet inspection," in *Proc. ACM/IEEE Symp. Archit. Netw. Commun. Syst.*, Dec. 2006, pp. 93–102.

[80] R. Sidhu and V. K. Prasanna, "Fast regular expression matching using FPGAs," in *Proc. Annu. IEEE Symp. Field-Program. Custom Comput. Mach.*, Mar./Apr. 2001, pp. 227–238.

[81] R. Antonello *et al.* "Deep packet inspection tools and techniques in commodity platforms: Challenges and trends," *J. Netw. Comput. Appl.*, vol. 35, no. 6, pp. 1863–1878, 2012.

[82] Intel. *Intel Pro/100S Network Adapters. IPSec Offload Performance and Comparison*. [Online]. Available: https://www.intel.com/content/dam/doc/performance-brief/etesting-labs-r%eport-intel-pro-100-s-network-adapters-brief.pdf

[83] Q. Zhang, V. Liu, H. Zeng, and A. Krishnamurthy, "High-resolution measurement of data center microbursts," in *Proc. Internet Meas. Conf.*, 2017, pp. 78–85.

[84] S. Pontarelli, G. Bianchi, and M. Welzl, "A programmable hardware calendar for high resolution pacing," in *Proc. IEEE 19th Int. Conf. High Perform. Switching Routing (HPSR)*, 2018.

[85] A. Caulfield, P. Costa, and M. Ghobadi, "Beyond smartnics: Towards a fully programmable cloud," in *Proc. IEEE Int. Conf. High Perform. Switching Routing*, 2018, pp. 1–6.

[86] *Cavium LiquidIO II Network Appliance Smart NICs*. [Online]. Available: http://www.cavium.com/LiquidIO-II_Network_Appliance_Adapters.html

[87] *BlueField Multicore System on Chip*. [Online]. Available: http://www.mellanox.com/related-docs/npu-multicore-processors/PB_Bluef%ield_SoC.pdf

[88] *Open vSwitch Offload and Acceleration With Agilio CX SmartNICs*. [Online]. Available: https://www.netronome.com/media/redactor_files/WP_OVS_Benchmarking.p%df

[89] *BlueField Multicore System on Chip*. [Online]. Available: http://www.mellanox.com/related-docs/prod_multi_core/PB_TILE-Gx72.pd%f

[90] D. Firestone *et al.*, "Azure accelerated networking: SmartNICs in the public cloud," in *Proc. USENIX Symp. Netw. Syst. Design Implement. (NSDI)*, 2018.

[91] N. McKeown *et al.*, "OpenFlow: Enabling innovation in campus networks," *ACM SIGCOMM Comput. Commun. Rev.*, vol. 38, no. 2, pp. 69–74, 2008.

[92] Intel. *Intel FlowDirector*. [Online]. Available: https://software.intel.com/en-us/articles/setting-up-intel-ethernet-flo%w-director

[93] P. Bosshart *et al.*, "Forwarding metamorphosis: Fast programmable match-action processing in hardware for SDN," in *ACM SIGCOMM Comput. Commun. Rev.*, vol. 43, no. 4, pp. 99–110, 2013.

[94] R. Duncan and P. Jungck, "packetC language for high performance packet processing," in *Proc. IEEE Int. Conf. High Perform. Comput. Commun. (HPCC)*, Jun. 2009, pp. 450–457.

[95] H. Song, "Protocol-oblivious forwarding: Unleash the power of SDN through a future-proof forwarding plane," in *Proc. ACM SIGCOMM Workshop Hot Topics Softw. Defined Netw.*, 2013, pp. 127–132.

[96] *Linux Socket Filtering AKA Berkeley Packet Filter (BPF)*. [Online]. Available: https://www.kernel.org/doc/Documentation/networking/filter.txt

[97] A. Sivaraman *et al.*, "Packet transactions: High-level programming for line-rate switches," in *Proc. ACM SIGCOMM Conf.*, 2016, pp. 15–28.

[98] G. Bianchi, M. Bonola, A. Capone, and C. Cascone, "OpenState: Programming platform-independent stateful openflow applications inside the switch," *ACM SIGCOMM Comput. Commun. Rev.*, vol. 44, no. 2, pp. 44–51, 2014.

[99] M. Moshref, A. Bhargava, A. Gupta, M. Yu, and R. Govindan, "Flow-level state transition as a new switch primitive for SDN," in *Proc. Workshop Hot Topics Softw. Defined Netw.*, 2014, pp. 61–66.

[100] G. Bianchi, M. Bonola, S. Pontarelli, D. Sanvito, A. Capone, and C. Cascone. (2016). "Open packet processor: A programmable architecture for wire speed platform-independent stateful in-network processing." [Online]. Available: https://arxiv.org/abs/1605.01977

[101] A. Caulfield, P. Costa, and M. Ghobadi, "Beyond SmartNICs: Towards a fully programmable cloud," in *Proc. IEEE HPSR*, 2018.

[102] M.-A. Kourtis *et al.*, "Enhancing VNF performance by exploiting SR-IOV and DPDK packet processing acceleration," in *Proc. IEEE Conf. Netw. Function Virtualization Softw. Defined Netw. (NFV-SDN)*, Nov. 2015, pp. 74–78.

[103] A. Panda, S. Han, K. Jang, M. Walls, S. Ratnasamy, and S. Shenker, "NetBricks: Taking the V out of NFV," in *Proc. USENIX OSDI*, 2016, pp. 203–216.

[104] S. Lange *et al.*, "Performance benchmarking of a software-based LTE SGW," in *Proc. Int. Workshop Manage. SDN NFV Syst. (ManSDN/NFV)*, Nov. 2015, pp. 378–383.

ABOUT THE AUTHORS

Leonardo Linguaglossa received the master's degree in telecommunication engineering from the University of Catania, Catania, Italy, in 2012 and the Ph.D. degree in computer networks, through a joint doctoral program, from Nokia Bell Labs, Nozay, France, Institut National de Recherche en Informatique et en Automatique, Paris, France, and University Paris 7, Paris, in 2016.

He is currently a Postdoctoral Researcher at Télécom ParisTech, Paris. His current research interests include architecture, design, and prototyping of systems for high-speed software packet processing, future Internet architecture, and software-defined networking.

Stanislav Lange received the M.Sc. degree in computer science from the University of Würzburg, Würzburg, Germany, in 2014, where he is currently working toward the Ph.D. degree.

He is currently a Researcher with the Next Generation Networks Research Group, Chair of Communication Networks, University of Würzburg. His current research interests include software-defined networking, performance analysis, system modeling, and multiobjective optimization.

Salvatore Pontarelli received the master's degree in electronic engineering from the University of Bologna, Bologna, Italy, and the Ph.D. degree in microelectronics and telecommunications from the University of Rome Tor Vergata, Rome, Italy.

He is currently a Senior Researcher at the Italian National Inter-University Consortium for Telecommunications, University of Rome Tor Vergata. His current research interests include hash-based structures for networking applications, use of field-programmable gate arrays for high-speed network monitoring, and hardware design of software-defined network devices.

Gábor Rétvári received the M.Sc. and Ph.D. degrees in electrical engineering from the Budapest University of Technology and Economics, Budapest, Hungary, in 1999 and 2007, respectively.

He is currently a Senior Research Fellow at the Department of Telecommunications and Media Informatics, Budapest University of Technology and Economics. He maintains several open-source scientific tools written in Perl, C, and Haskell. His current research interests include all aspects of network routing and switching, the programmable data plane, and the networking applications of computational geometry and information theory.

Dario Rossi (Senior Member, IEEE) received the M.Sc. and Ph.D. degrees from the Politecnico di Torino, Turin, Italy, in 2001 and 2005, respectively, and the HDR degree from the Université Pierre et Marie Curie, Paris, France, in 2010.

He is currently a Professor at Télécom ParisTech, Paris, and École Polytechnique, Palaiseau, France. He has coauthored nine patents and over 150 papers in leading conferences and journals.

Dr. Rossi is a Senior Member of the Association for Computing Machinery. He holds Cisco's Chair NewNet@Paris. He received eight best paper awards, the Google Faculty Research Award in 2015, and the IRTF Applied Network Research Prize in 2016.

Thomas Zinner received the Ph.D. and Habilitation degrees from the University of Würzburg, Würzburg, Germany, in 2012 and 2017, respectively.

Since 2018, he has been a Visiting Professor and the Head of the Internet Network Architectures Research Group, Technical University of Berlin, Berlin, Germany. He has published more than 80 research papers in major conferences and journals, six of which received best paper and best student paper awards. His current research interests include quality of experience management, network softwarization, programmable data planes, and performance evaluation.

Roberto Bifulco received the Ph.D. degree from the University of Naples Federico II, Naples, Italy.

He is currently a Senior Researcher with the Systems and Machine Learning Group, NEC Laboratories Europe, Heidelberg, Germany. Before joining NEC in 2012, he was a consultant for small technological enterprises and start-ups in the fields of server virtualization and cloud computing. His current research interests include the design of high-performance networked systems and the exploration of novel solutions at the crossroad of systems and machine learning.

Michael Jarschel received the Ph.D. degree from the University of Würzburg, Würzburg, Germany, in 2014, with the thesis titled "An Assessment of Applications and Performance Analysis of Software Defined Networking."

He is currently a Research Engineer in the areas of network softwarization and connected driving at Nokia Bell Labs, Munich, Germany. His current research interests include the applicability of software-defined networking and network function virtualization concepts to next-generation networks and distributed telco cloud technologies and their use cases.

Giuseppe Bianchi is currently a Full Professor of Networking at the University of Rome Tor Vergata, Rome, Italy. His current research interests include programmable network systems, wireless networks, privacy and security, and traffic control.

Dr. Bianchi has been the general or the technical co-chair for several major IEEE and Association for Computing Machinery (ACM) conferences and workshops. So far, he has held general or technical coordination roles in six European projects. He has been an Editor of the IEEE/ACM TRANSACTIONS ON NETWORKING, the IEEE TRANSACTIONS ON WIRELESS COMMUNICATIONS, the IEEE TRANSACTIONS ON NETWORK AND SERVICE MANAGEMENT, and *Computer Communications* (Elsevier).