

# Overview of last class

# Cookies: keeping “state” (cont.)

client



server



cookie file



ebay 8734  
amazon 1678

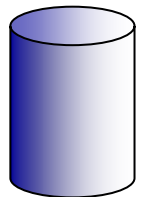
usual http request msg

Amazon server  
creates ID  
1678 for user

usual http response  
**set-cookie: 1678**

create  
entry

backend  
database



usual http request msg  
**cookie: 1678**

cookie-  
specific  
action

access

usual http response msg

access

cookie-  
specific  
action

one week later:



ebay 8734  
amazon 1678

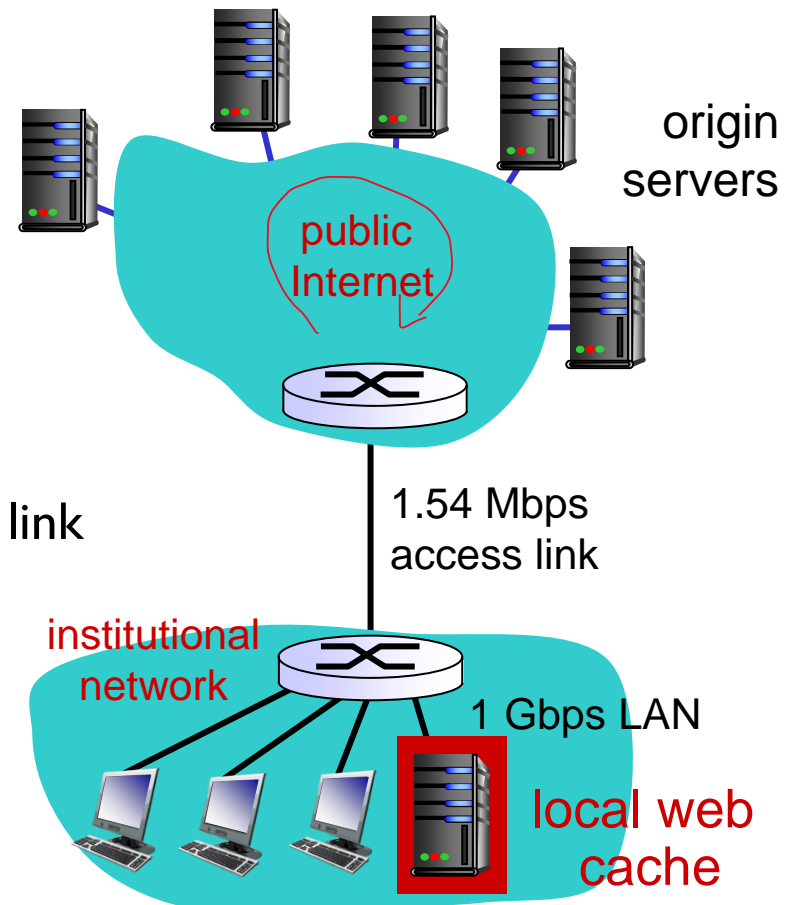
usual http request msg  
**cookie: 1678**

usual http response msg

# Caching example: install local cache

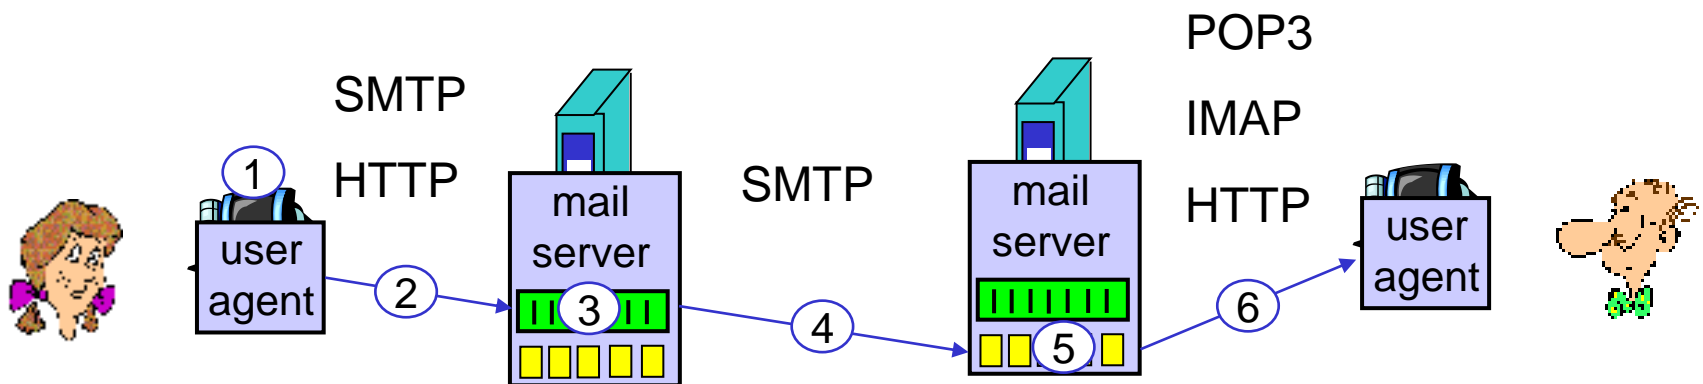
## *Calculating access link utilization, delay with cache:*

- suppose cache hit rate is 0.4
  - 40% requests satisfied at cache, 60% requests satisfied at origin
- access link utilization:
  - 60% of requests use access link
- data rate to browsers over access link
  - $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
  - utilization =  $0.9 / 1.54 = .58$
- total delay
  - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
  - $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$
  - less than with 154 Mbps link (and cheaper too!)



# Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message and “to”  
`bob@someschool.edu`
- 2) Alice’s UA sends message to her mail server; message placed in message queue
- 3) Client side of SMTP opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



# SMTP: final words

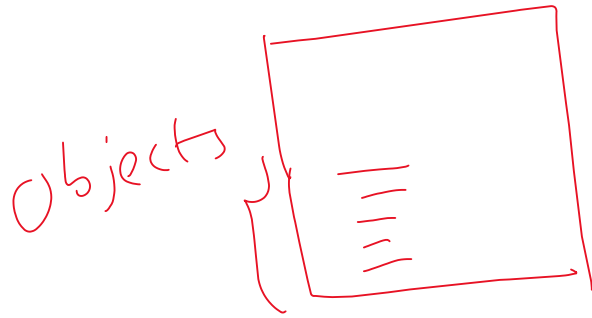
- SMTP uses persistent connections
- SMTP requires message (header & body) to be in 7-bit ASCII
- SMTP server uses CRLF.CRLF to determine end of message
- The MIME (multipurpose Internet mail extension) extensions for non-ascii data

Q: why persistent mode is selected?

## *comparison with HTTP:*

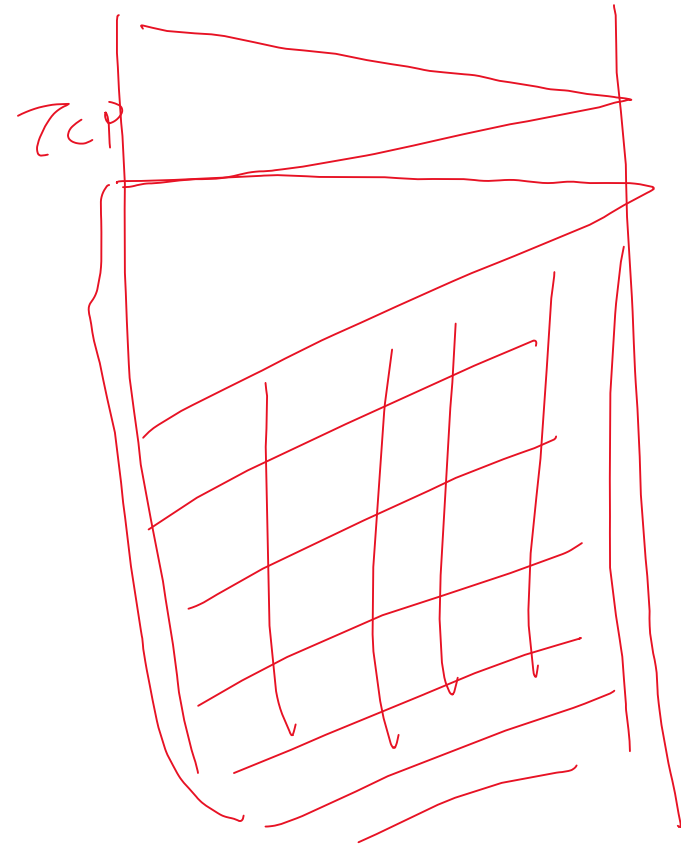
- HTTP: pull
- SMTP: push
- both have ASCII command/response interaction, status codes
- HTTP: each object encapsulated in its own response message
- SMTP: multiple objects sent in multipart message

Alice → Bob



Each objects  
normally small size

- persistent: push all the objects one by one
- non-persistent
  - o non-parallel → bad choice
  - o parallel → collection management congestion/flow control incur cost



# Class Today

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP



# DNS: domain name system

*people*: many identifiers:

- SSN, name, passport #

*Internet hosts, routers*:

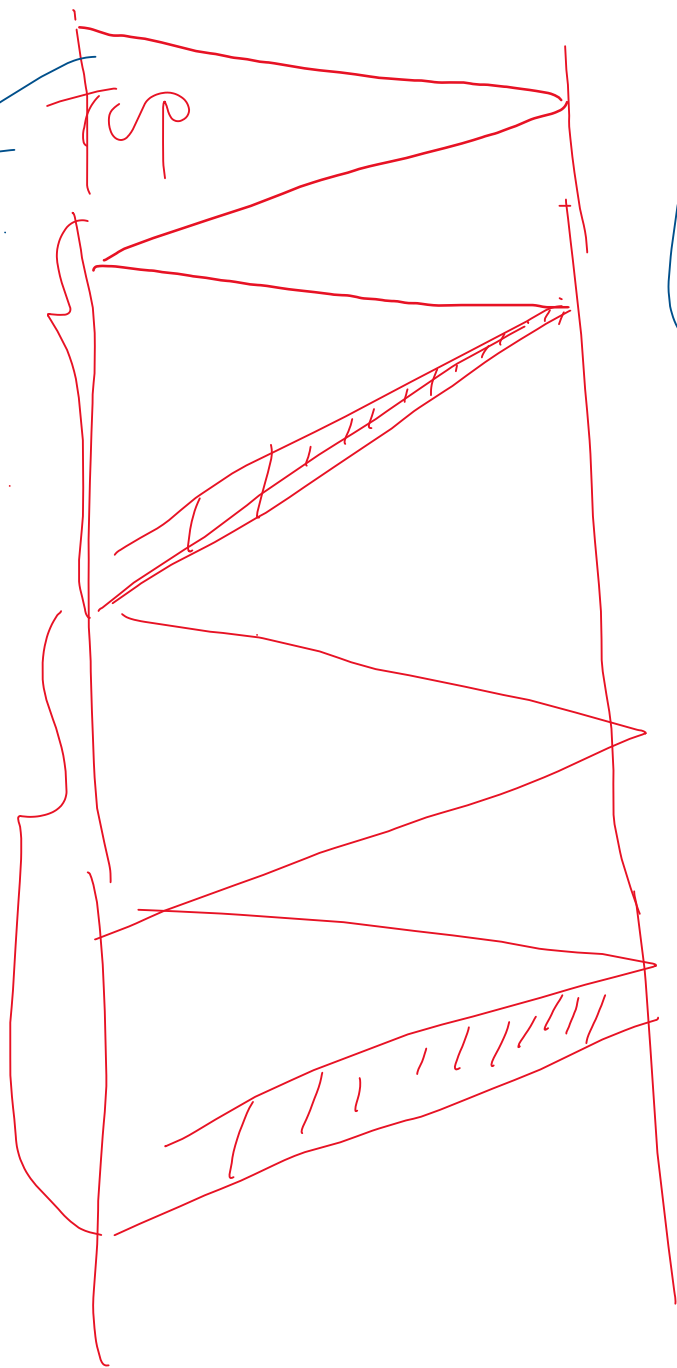
- IP address (32 bit) - used for addressing datagrams
- “name”, e.g.,  
www.yahoo.com - used by humans

Q: how to map between IP address and name?

## *Domain Name System:*

- *distributed database*  
implemented in hierarchy of many *name servers*
- *application-layer protocol*: hosts, name servers communicate to *resolve* names (address/name translation)
  - note: **core Internet function, implemented as application-layer protocol**
  - complexity at network's “edge”
- *Runs over UDP: why not TCP?*

state  
many

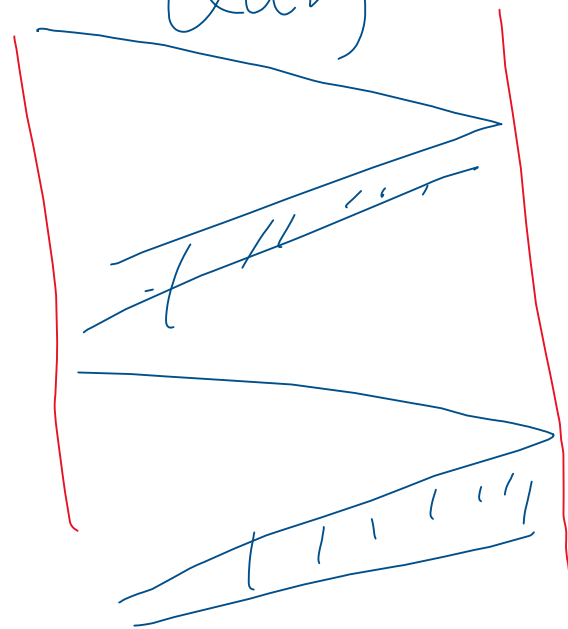


Scalability

Delay

Band.  
Efficiency

Queue



# DNS: services, structure

## *DNS services*

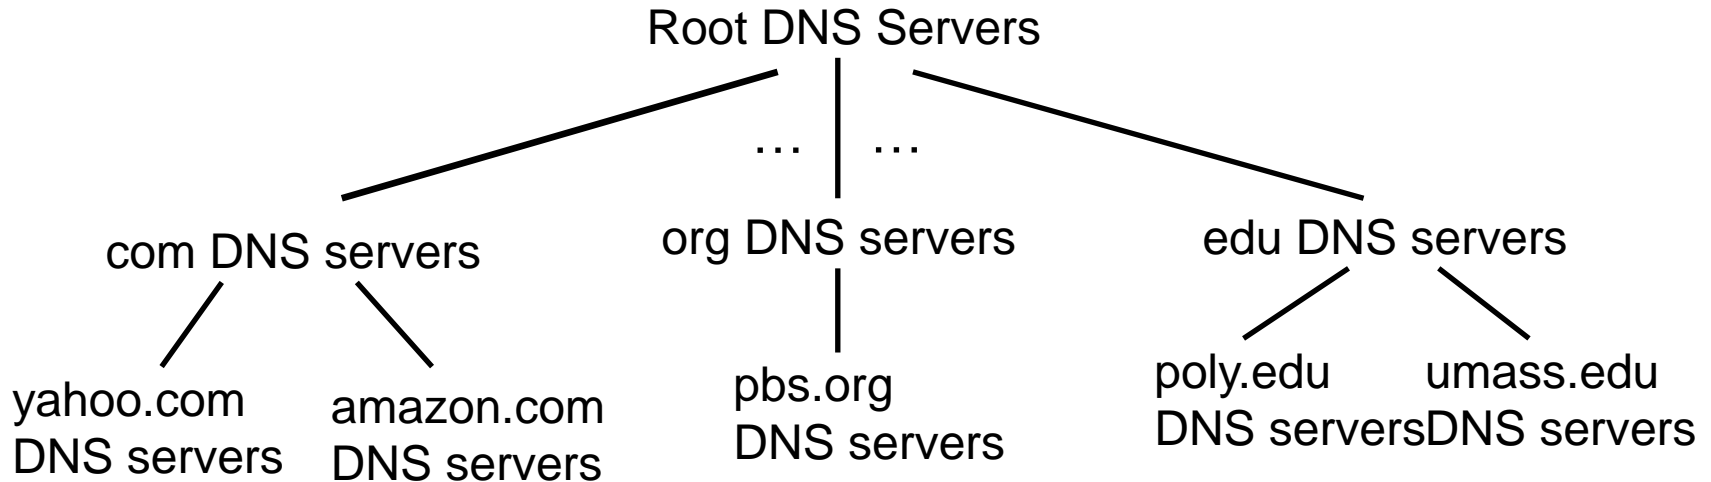
- hostname to IP address translation
- host aliasing
  - canonical, alias names
- mail server aliasing
- load distribution
  - replicated Web servers: many IP addresses correspond to one name

## *why not centralize DNS?*

- single point of failure
- traffic volume
- distant centralized database
- maintenance

*A: doesn't scale!*

# DNS: a distributed, hierarchical database

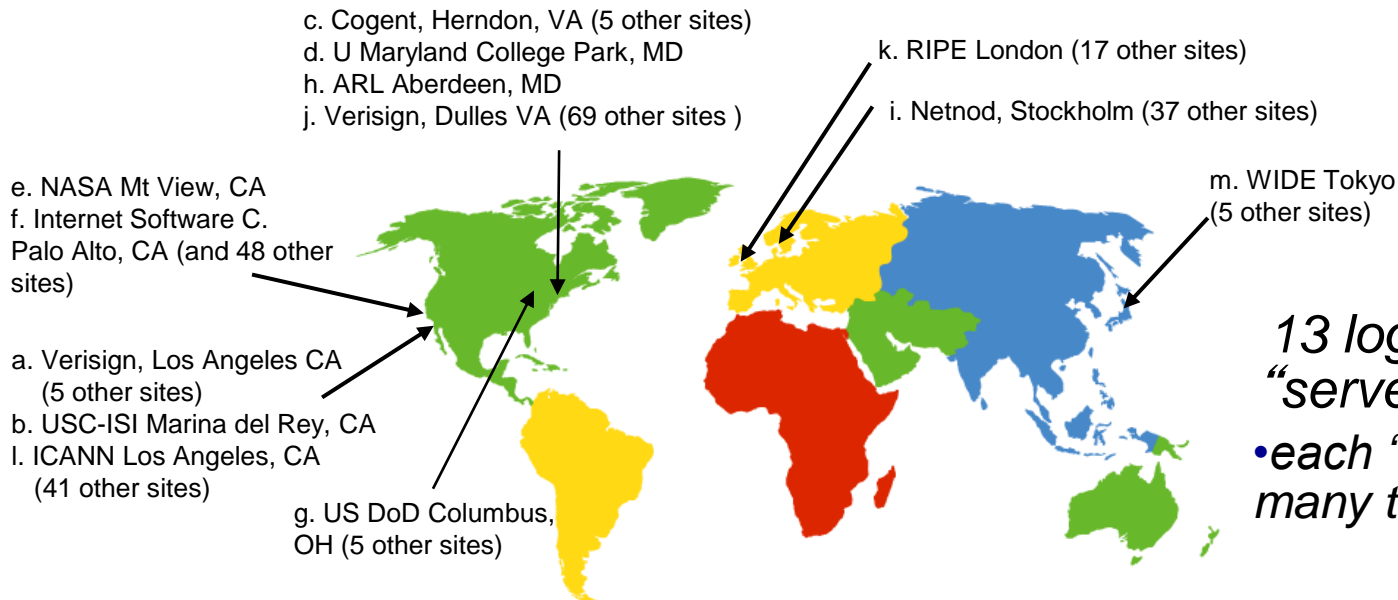


*client wants IP for [www.amazon.com](http://www.amazon.com):*

- client queries root server to find com DNS server
- client queries .com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for [www.amazon.com](http://www.amazon.com)

# DNS: root name servers

- contacted by local name server that can not resolve name
- root name server:
  - contacts authoritative name server if name mapping not known
  - gets mapping
  - returns mapping to local name server



*13 logical root name  
“servers” worldwide*  
• *each “server” replicated  
many times*

# TLD, authoritative servers

## *top-level domain (TLD) servers:*

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
- Network Solutions maintains servers for .com TLD
- Educause for .edu TLD

## *authoritative DNS servers:*

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

# Local DNS name server

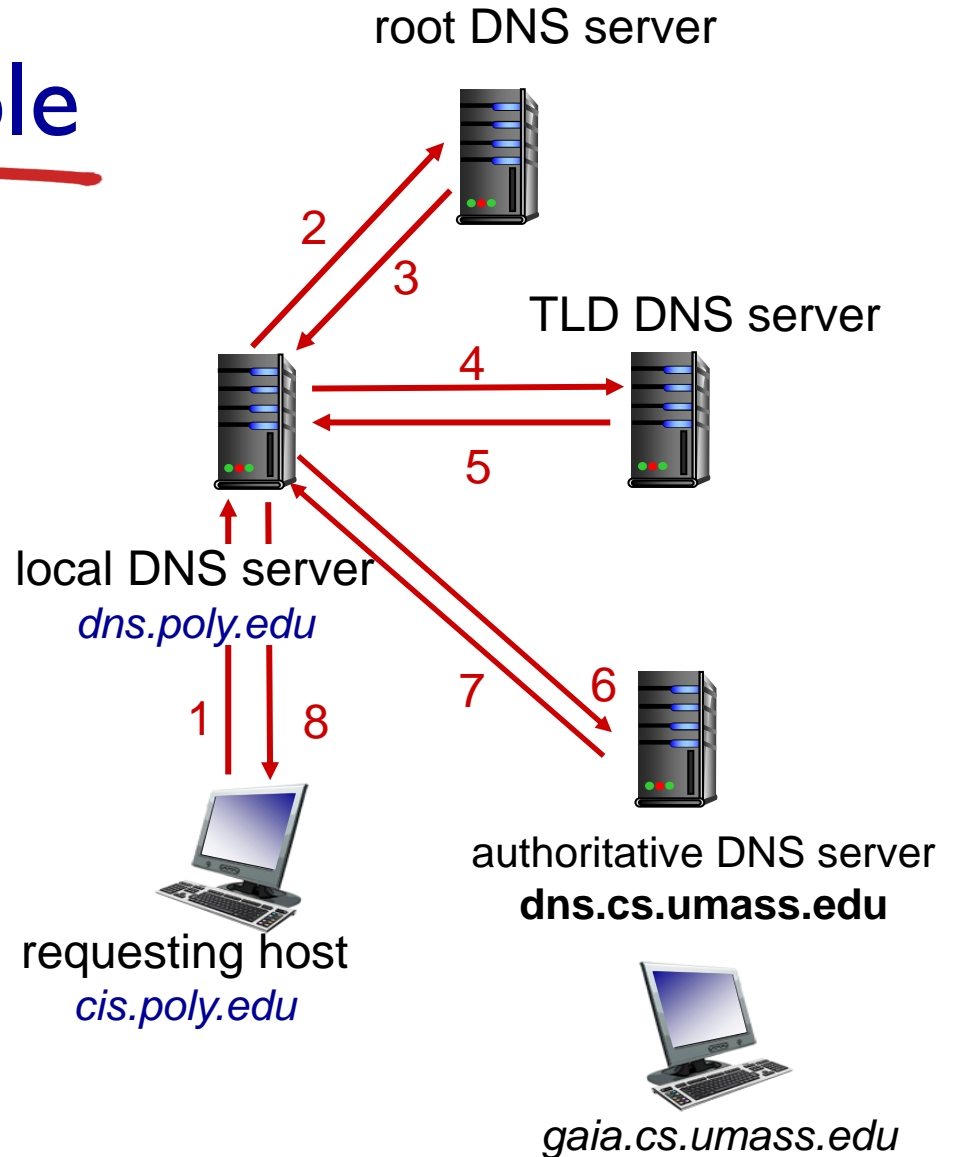
- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one
  - also called “default name server”
- when host makes DNS query, query is sent to its local DNS server
  - has local cache of recent name-to-address translation pairs (but may be out of date!)
  - acts as proxy, forwards query into hierarchy

# DNS name resolution example

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

## *iterated query:*

- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”

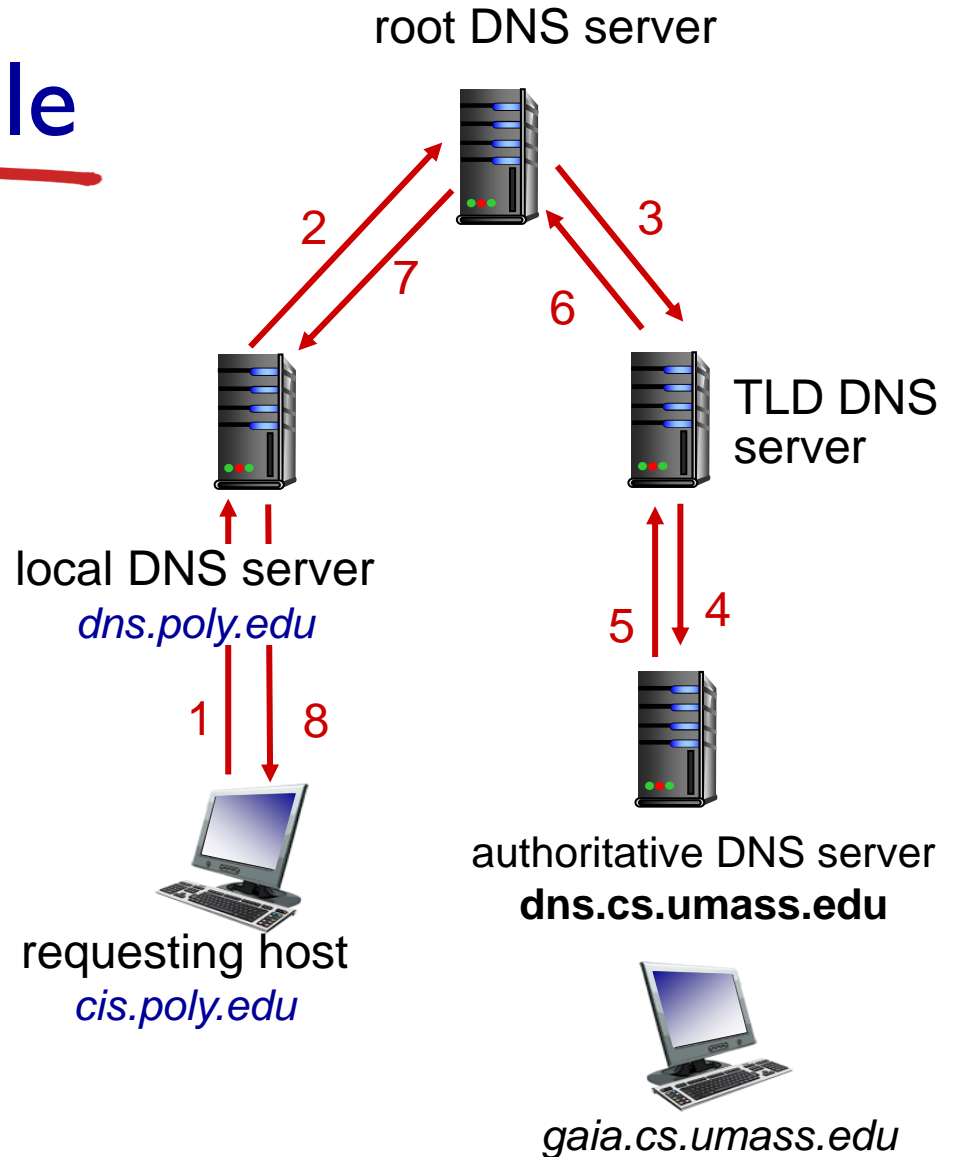




# DNS name resolution example

## *recursive query:*

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy?



# DNS: caching, updating records

- once (any) name server learns mapping, it *caches* mapping
  - cache entries timeout (disappear) after some time (TTL)
  - TLD servers typically cached in local name servers
    - thus root name servers not often visited
- cached entries may be *out-of-date* (best effort name-to-address translation!)
  - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- update/notify mechanisms proposed IETF standard
  - RFC 2136

# DNS records

**DNS:** distributed database storing resource records (RR)

RR format: (name, value, type, ttl)

## type=A

- **name** is hostname
- **value** is IP address

## type=NS

- **name** is domain (e.g., foo.com)
- **value** is hostname of authoritative name server for this domain

## type=CNAME

- **name** is alias name for some “canonical” (the real) name
- **www.ibm.com** is really **servereast.backup2.ibm.com**
- **value** is canonical name

## type=MX

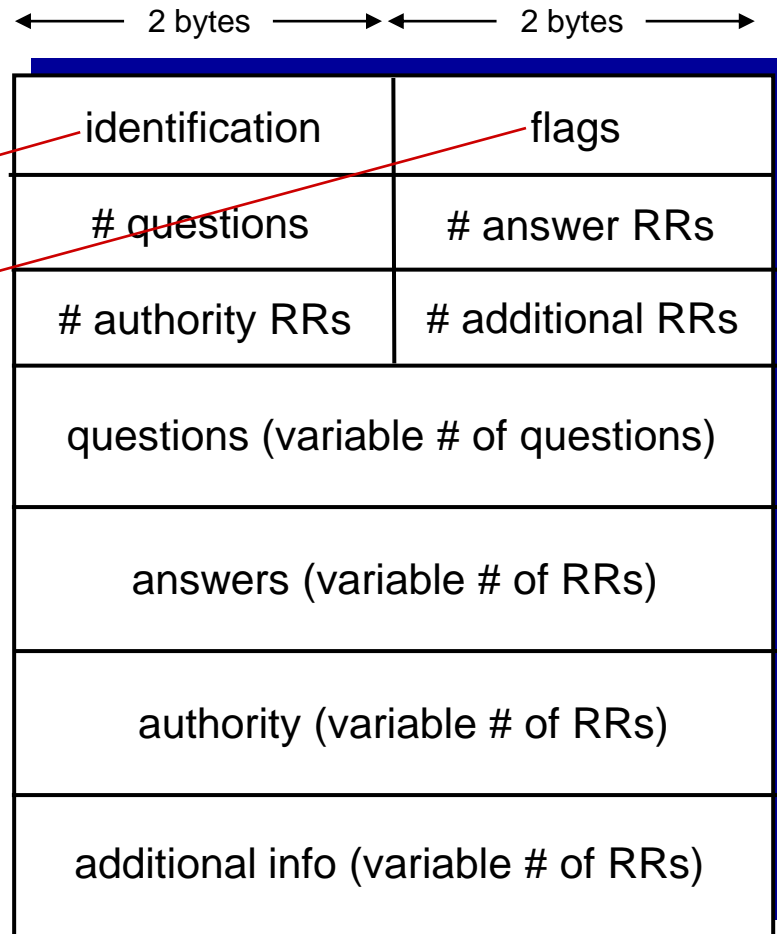
- **value** is name of mailserver associated with **name**

# DNS protocol, messages

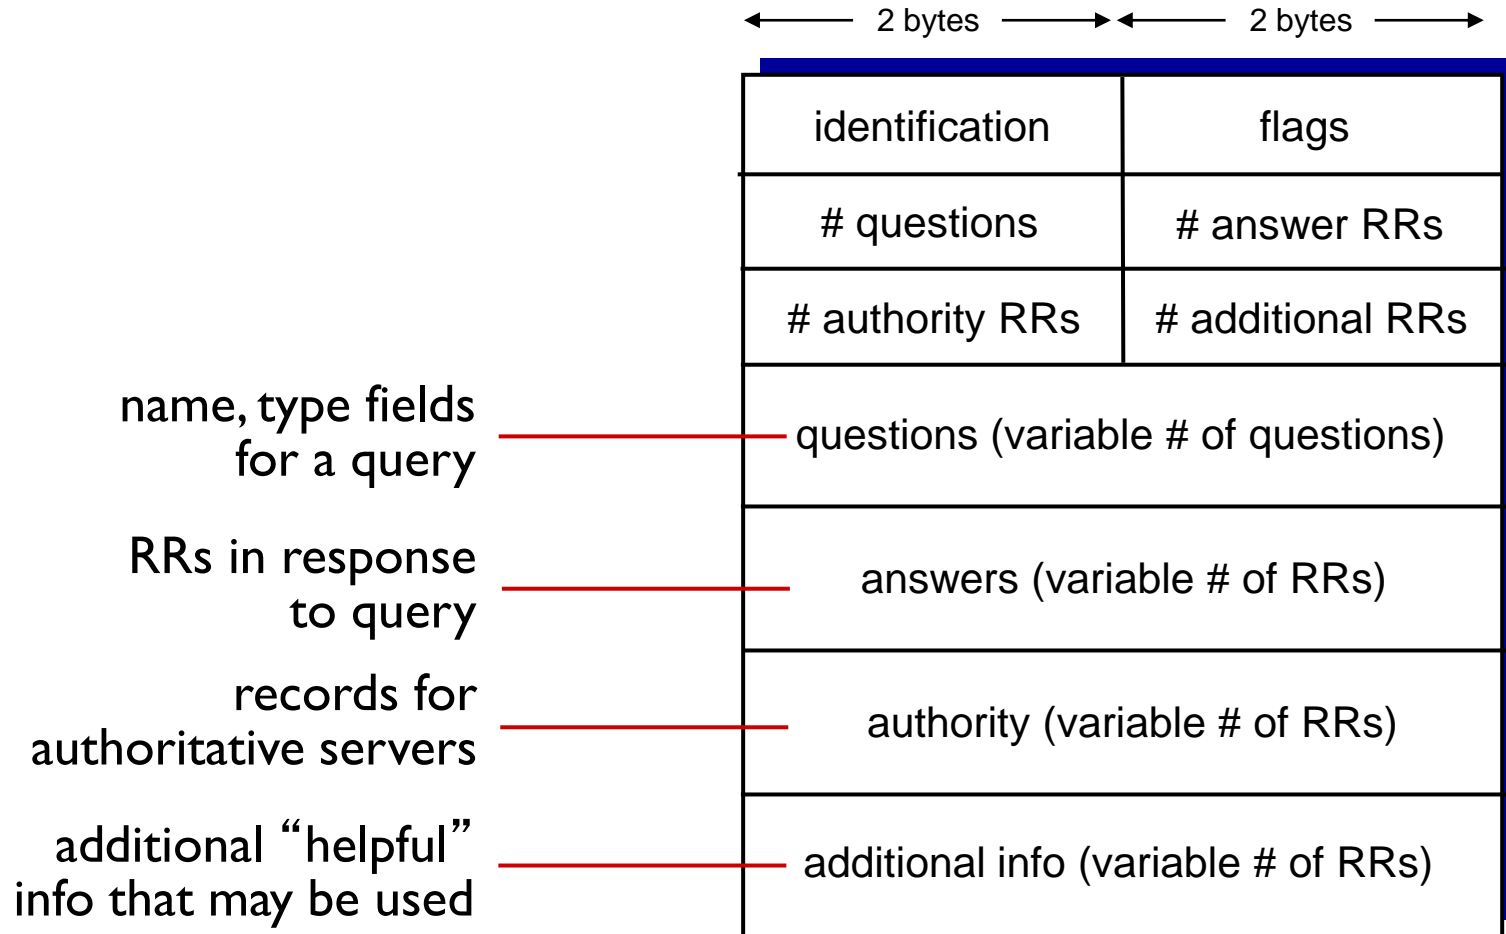
- *query* and *reply* messages, both with same *message format*

## message header

- **identification:** 16 bit # for query, reply to query uses same #
- **flags:**
  - query or reply
  - recursion desired
  - recursion available
  - reply is authoritative



# DNS protocol, messages



### Header flags format

| Field  | Description  | Length (bits) |
|--------|--|---------------|
| QR     | Indicates if the message is a query (0) or a reply (1)   | 1             |
| OPCODE | The type can be QUERY (standard query, 0), IQUERY (inverse query, 1), or STATUS (server status request, 2)                         | 4             |
| AA     | Authoritative Answer, in a response, indicates if the DNS server is authoritative for the queried hostname                         | 1             |
| TC     | TrunCation, indicates that this message was truncated due to excessive length  | 1             |
| RD     | Recursion Desired, indicates if the client means a recursive query   | 1             |
| RA     | Recursion Available, in a response, indicates if the replying DNS server supports recursion  | 1             |
| Z      | Zero, reserved for future use  | 3             |
| RCODE  | Response code, can be NOERROR (0), FORMERR (1, Format error), SERVFAIL (2), NXDOMAIN (3, Nonexistent domain), etc. <sup>[34]</sup> | 4             |

[https://en.wikipedia.org/wiki/Domain\\_Name\\_System](https://en.wikipedia.org/wiki/Domain_Name_System)

- ▼ Domain Name System (query)
  - Transaction ID: 0x178e
  - > Flags: 0x0100 Standard query
  - Questions: 1
  - Answer RRs: 0
  - Authority RRs: 0
  - Additional RRs: 0
  - ▼ Queries
    - ▼ image.google.com: type A, class IN

Name: image.google.com

[Name Length: 16]

[Label Count: 3]

Type: A (Host Address) (1)

Class: IN (0x0001)

[Response In: 7]

|      |   |                  |
|------|---|------------------|
| 0000 | 02 00 00 00 00 04 3c 22 fb c2 d2 51 08 00 45 00 | .....<" ...Q..E. |
| 0010 | 00 3e 61 1d 00 00 40 11 ab c4 c0 a8 00 05 d0 43 | ·>a···@· .....C  |
| 0020 | dc dc de 06 00 35 00 2a ae 19 17 8e 01 00 00 01 | .....5·* .....   |
| 0030 | 00 00 00 00 00 00 05 69 6d 61 67 65 06 67 6f 6f | .....·i mage·goo |
| 0040 | 67 6c 65 03 63 6f 6d 00 00 01 00 01             | gle·com· ....    |

<https://cabulous.medium.com/dns-message-how-to-read-query-and-response-message-cfebcb4fe817>

## Domain Name System (response)

Transaction ID: 0x178e

> Flags: 0x8180 Standard query response, No error

Questions: 1

Answer RRs: 3

Authority RRs: 0

Additional RRs: 0

> Queries

Answers

> image.google.com: type CNAME, class IN, cname images.google.com

> images.google.com: type CNAME, class IN, cname images.l.google.com

> images.l.google.com: type A, class IN, addr 172.217.1.14

[Request In: 5]

[Time: 0.038051000 seconds]

|      |   |                   |
|------|---|-------------------|
| 0000 | 3c 22 fb c2 d2 51 02 00 00 00 00 04 08 00 45 00 | <"...Q.. ....E.   |
| 0010 | 00 7a 15 b0 40 00 3b 11 bb f5 d0 43 dc dc c0 a8 | .z..@.;. ...C.... |
| 0020 | 00 05 00 35 de 06 00 66 e1 b0 17 8e 81 80 00 01 | ...5...f .....    |
| 0030 | 00 03 00 00 00 00 05 69 6d 61 67 65 06 67 6f 6f | .....i mage·goo   |
| 0040 | 67 6c 65 03 63 6f 6d 00 00 01 00 01 c0 0c 00 05 | gle·com· .....    |
| 0050 | 00 01 00 00 00 3c 00 09 06 69 6d 61 67 65 73 c0 | .....<.. ·images· |
| 0060 | 12 c0 2e 00 05 00 01 00 09 3a 80 00 0b 06 69 6d | ..,..... :.....im |
| 0070 | 61 67 65 73 01 6c c0 12 c0 43 00 01 00 01 00 00 | ages·l·· ·C·..... |
| 0080 | 01 2c 00 04 ac d9 01 0e                         | ·,.....           |



# Inserting records into DNS

- example: new startup “Network Utopia”
- register name networkutopia.com at *DNS registrar* (e.g., Network Solutions)
  - provide names, IP addresses of authoritative name server (primary and secondary)
  - registrar inserts two RRs into .com TLD server:  
(networkutopia.com, dns1.networkutopia.com, NS)  
(dns1.networkutopia.com, 212.212.212.1, A)
- create authoritative server type A record for www.networkutopia.com; type MX record for networkutopia.com and type A record for the email server

# Attacking DNS

## DDoS attacks

- bombard root servers with traffic
  - not successful to date
  - traffic filtering
  - local DNS servers cache IPs of TLD servers, allowing root server bypass
- bombard TLD servers
  - potentially more dangerous

## redirect attacks

- man-in-middle
  - Intercept queries
- DNS poisoning
  - Send bogus replies to DNS server, which caches

## exploit DNS for DDoS

- send queries with spoofed source address: target IP
- requires amplification